

Lecture #13


Complex Data and Control Structures (2): structs and their alignment

Section 3.9

Prof. Soraya Abad-Mota, PhD

Structure Representation


```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- ▶ Structure represented as block of memory
 - **Big enough to hold all of the fields**
- ▶ Fields ordered according to declaration
 - **Even if another ordering could yield a more compact representation**
- ▶ Compiler determines overall size + positions of fields
 - **Machine-level program has no understanding of the structures in the source code**

Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- ▶ Generating Pointer to Array Element
 - Offset of each structure member determined at compile time
 - Compute as $r + 4 * idx$

```
int *get_ap  
(struct rec *r, size_t idx)  
{  
    return &r->a[idx];  
}
```


```
# r in %rdi, idx in %rsi  
leaq (%rdi,%rsi,4), %rax  
ret
```

Following Linked List

► C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    int i;
    struct rec *next;
};
```



Register	Value
%rdi	r
%rsi	val

```
.L11:                                # loop:
    movslq    16(%rdi), %rax          # i = M[r+16]
    movl      %esi, (%rdi,%rax,4)     # M[r+4*i] = val
    movq      24(%rdi), %rdi          # r = M[r+24]
    testq     %rdi, %rdi              # Test r
    jne       .L11                    # if !=0 goto loop
```

Practice

- ▶ Problem 3.41 p. 268

```
struct prob {  
    int *p;  
    struct {  
        int x;  
        int y;  
    } s;  
    struct prob *next;  
};
```

Problem 3.41 (cont.)

A. Offsets (in bytes) of the following fields?

p:

s.x:

s.y:

next:

B. How many total bytes does the structure require?

Problem 3.41 (cont.)

C. Assembly code for void sp_init(struct prob *sp)

```
    movl 12(%rdi), %eax  
    movl %eax, 8(%rdi)  
    leaq 8(%rdi), %rax  
    movq %rax, (%rdi)  
    movq %rdi, 16(%rdi)  
    ret
```

Fill in the C code for sp_init

Skeleton code for sp_init


```
void sp_init(struct prob *sp) {  
    sp->s.x =  
    sp->p    =  
    sp->next =
```


Practice more with

- ▶ Problem 3.42
- ▶ You may do it on your own.
- ▶ If we have time today, we can go over it at the end, once we have finished structure alignment and unions, but please make sure you do it.

Structures & Alignment

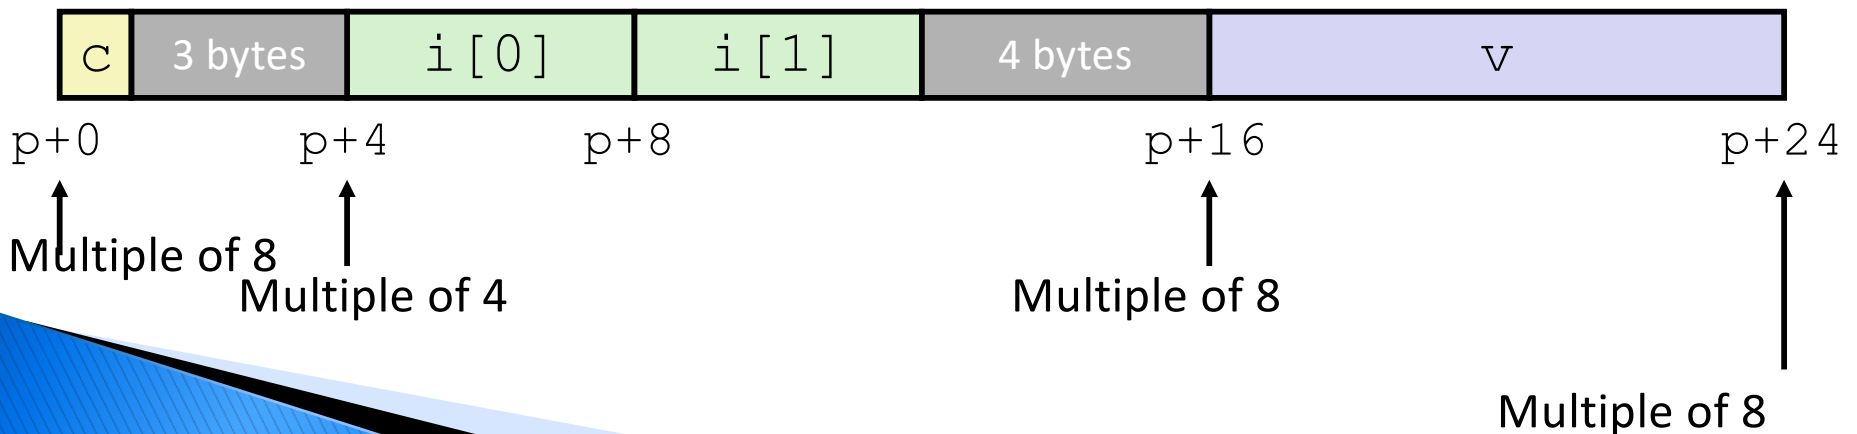
▶ Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

▶ Aligned Data

- Primitive data type requires K bytes
- (its) Address (inside the structure) must be multiple of K



Alignment Principles

▶ Aligned Data

- Primitive data type require K bytes → Address must be multiple of K
- Required on some machines; advised on x86-64

▶ Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (syst. dependent)
 - Inefficient to load or store datum that spans quad word boundaries
 - Virtual memory trickier when datum spans 2 pages

▶ Compiler

- Inserts gaps in structure to ensure correct *alignment of fields*

Specific Cases of Alignment (x86-64)

- ▶ K
- ▶ 1 byte: **char**, ...
 - no restrictions on address
- ▶ 2 bytes: **short**, ...
 - lowest 1 bit of address must be 0₂
- ▶ 4 bytes: **int**, **float**, ...
 - lowest 2 bits of address must be 00₂
- ▶ 8 bytes: **double**, `long`, **char ***, ...
 - lowest 3 bits of address must be 000₂
- ▶ 16 bytes: **long double** (GCC on Linux)
 - lowest 4 bits of address must be 0000₂

Satisfying Alignment with Structures

▶ Within structure:

- Must satisfy each **field's** alignment requirement

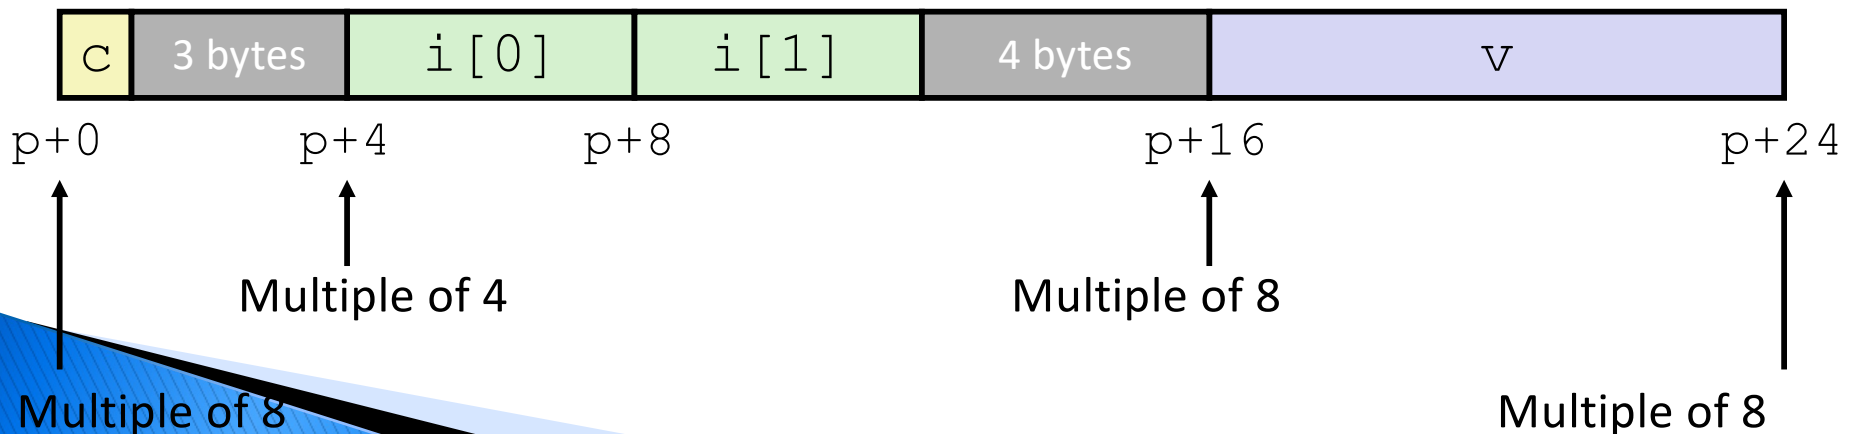
▶ Overall structure placement

- The **structure has alignment requirement K_{struct}**
 - K_{struct} = Largest alignment of any field
- **Initial address & structure length** must be multiples of K_{struct}

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

▶ Example:


- $K_{\text{struct}} = 8$, due to **double** element



Meeting Overall Alignment Requirement

- ▶ For largest alignment requirement K
- ▶ Overall structure must be multiple of K_{struct}
- ▶ More compact but requires user providing diff. field order

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```




Multiple of $K=8$

Practice Problem 3.44

- ▶ A. struct P1 {int l; char c; int j; char d;};
 - offset of each field?
 - total size of the structure?
 - alignment requirement, $K = ?$

Practice Problem 3.44 (cont.)


- ▶ B. `struct P2 {int i; char c; char d; long j;};`
 - offset of each field?
 - total size of the structure?
 - alignment requirement, $K = ?$
- ▶ Continue with parts C, D and E of this problem



Arrays of Structures

- ▶ Overall structure length multiple of K
- ▶ Must satisfy alignment requirement for every element (aligned to 24 bytes; where is the padding?)

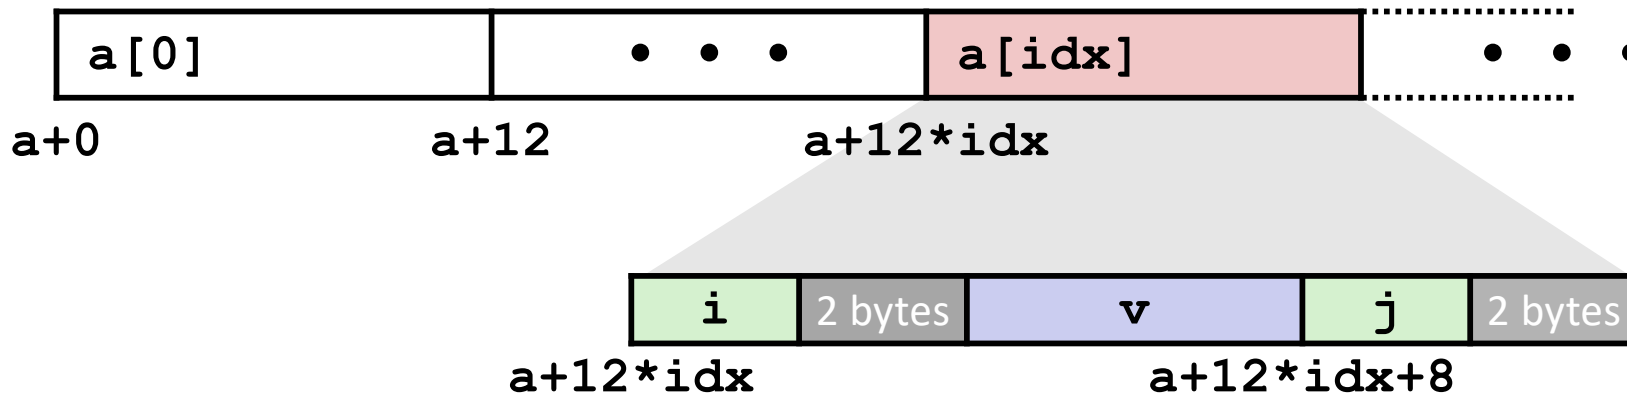
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Accessing Array Elements

- ▶ Compute array offset $12 * \text{idx}$
 - `sizeof(S3)`, including alignment spacers
- ▶ Element `j` is at offset 8 within structure
- ▶ Assembler gives offset `a+8`
 - Resolved during linking

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```




```
short get_j(int idx)  
{  
    return a[idx].j;  
}
```

```
# %rdi = idx  
leaq (%rdi,%rdi,2),%rax # 3*idx  
movzwl a+8(,%rax,4),%eax
```

Saving Space



- ▶ Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

- ▶ Effect (K=4) (Multiples of 4)



But this must be done by the programmer, the compiler will not change the order of the fields in the structure, it only adds padding.

End of section 3.9.1
and 3.9.3

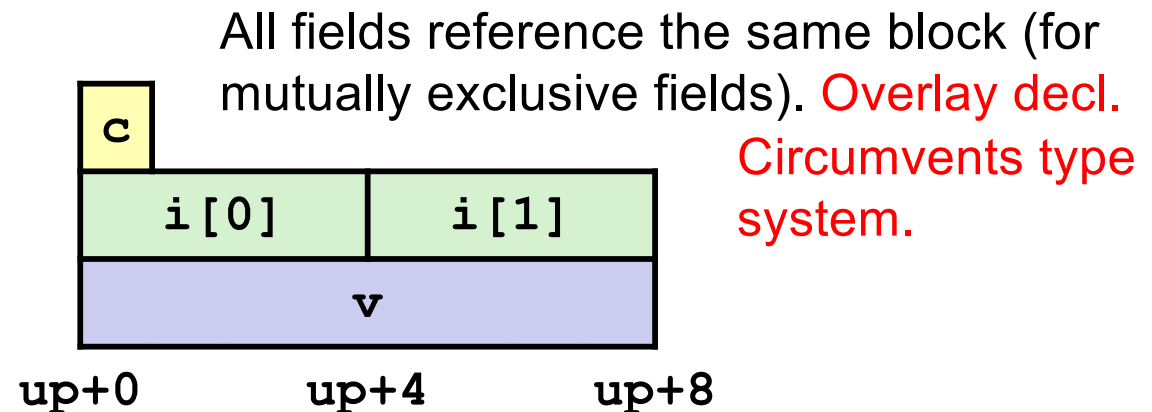
Now section 3.9.2 on
Unions

Union Allocation

- ▶ Allocate according to largest element
- ▶ Can only use one field at a time

```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```




Unlike struct in which diff. fields reference diff. blocks



Using Union to Access Bit Patterns

```
typedef union {  
    float f;  
    unsigned u;  
} bit_float_t;
```



Same location

```
float bit2float(unsigned u)  
{  
    bit_float_t arg;  
    arg.u = u;  
    return arg.f;  
}
```

Same as **(float) u**?

```
unsigned float2bit(float f)  
{  
    bit_float_t arg;  
    arg.f = f;  
    return arg.u;  
}
```

Same as **(unsigned) f**?

Two separate locations

Unions, Structures, and Bitfields

- ▶ Can even break up structures and unions into smaller pieces:

```
typedef struct floatbits {  
    unsigned int fraction:23;  
    unsigned int exponent:8;  
    unsigned int sign:1;  
} floatbits;  
  
typedef union {  
    float f;  
    floatbits b;  
} floatunion;
```

Stop here