

Lecture #26

Virtual Memory

Prof. Soraya Abad-Mota, PhD

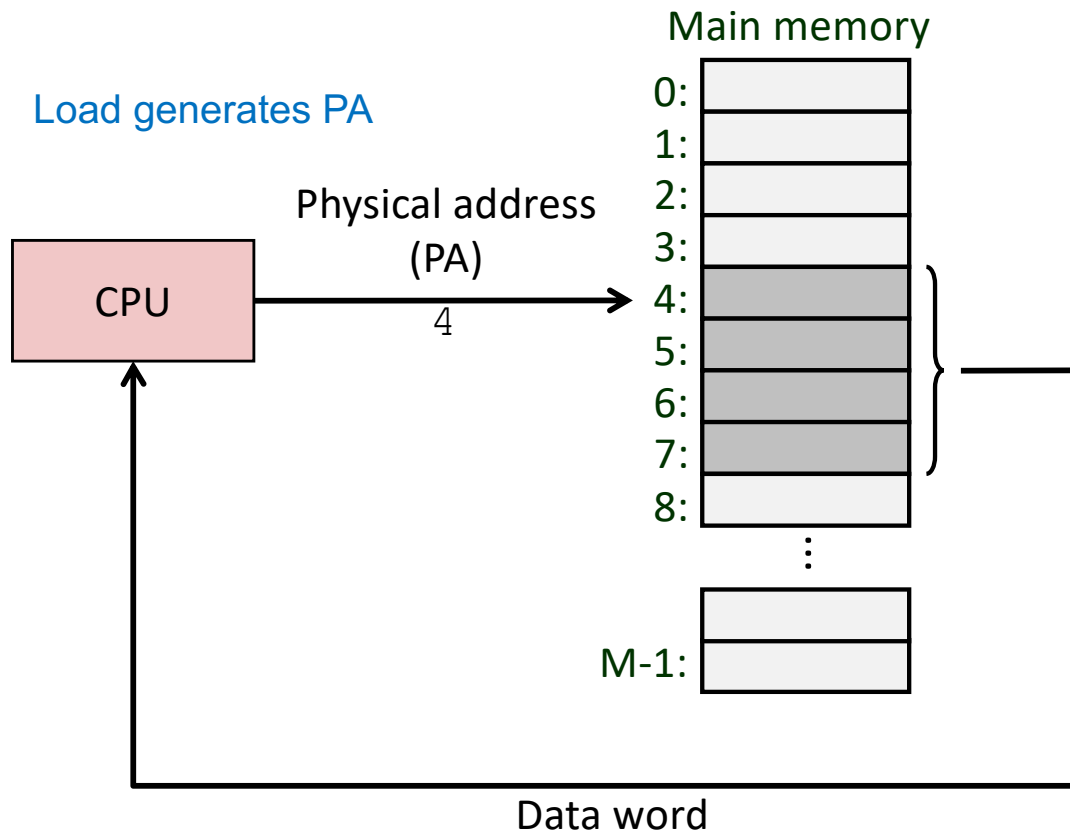
Topics

- ▶ **Virtual Memory** (material for the Final Exam)
 - Introduction
 - Address spaces
 - VM as a tool for caching
 - VM as a tool for memory management
 - VM as a tool for memory protection
 - Address translation (serves as review of cache)

Virtual Memory (VM)

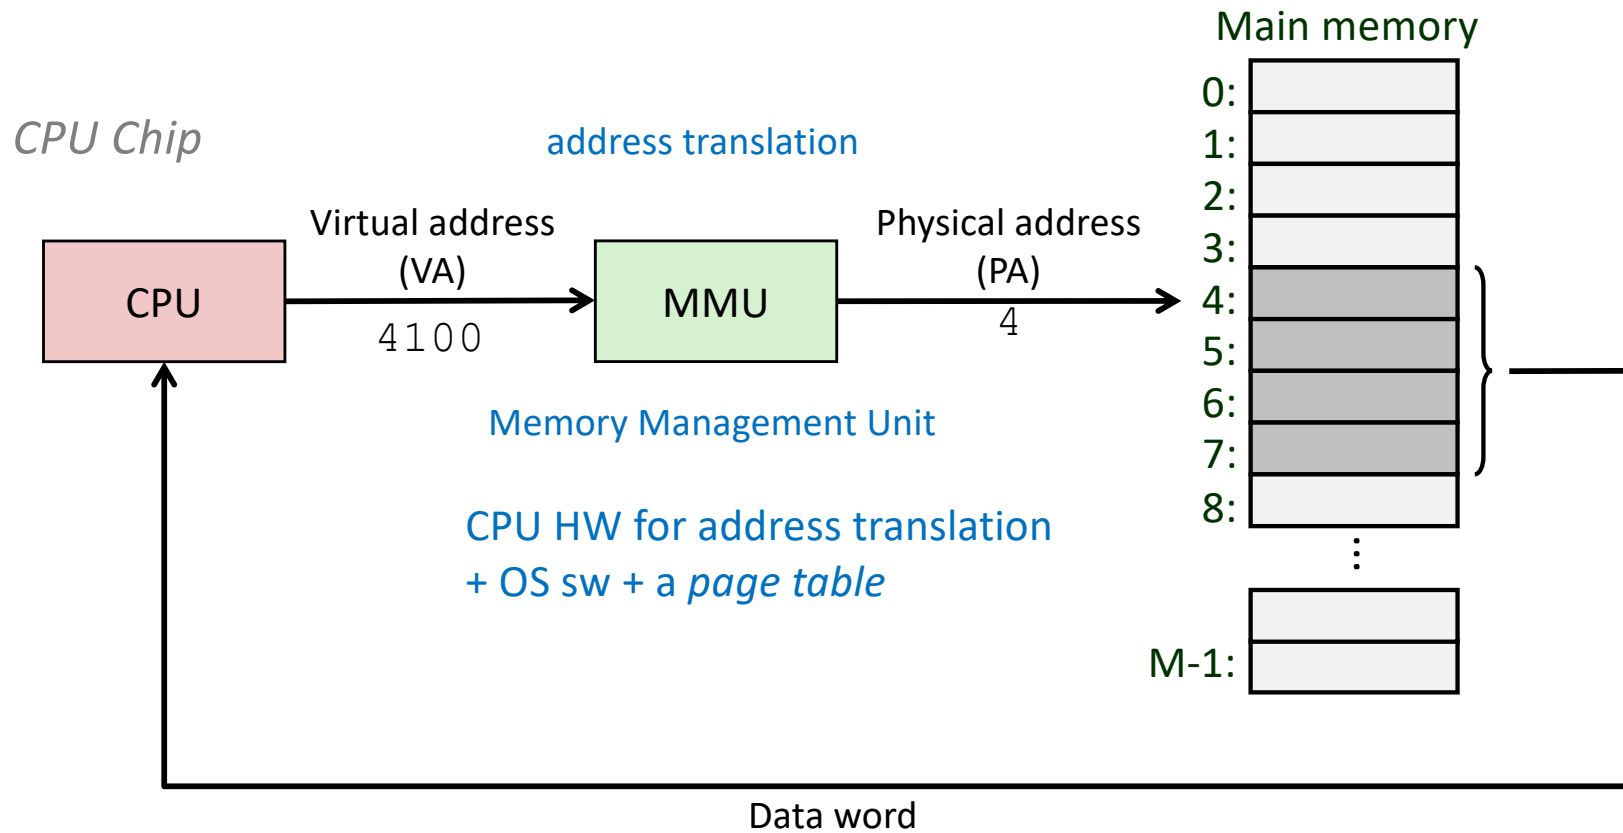
- ▶ One of the great ideas of computer systems
- ▶ Abstraction of main memory
- ▶ Designed to manage memory more efficiently and with fewer errors
- ▶ VM provides:
 - cache for address space stored in disk
 - to each process a uniform address space
 - protects the address space of each process from corruption by other processes.

A System Using Physical Addressing



- ▶ Used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and digital picture frames

A System Using Virtual Addressing



- ▶ Used in all modern servers, laptops, and smart phones
- ▶ One of the great ideas in computer science

Address Spaces

- ▶ **Linear address space:** Ordered set of contiguous non-negative integer addresses:

$\{0, 1, 2, 3 \dots \}$

- ▶ **Virtual address space:** Set of $N = 2^n$ virtual addresses

$\{0, 1, 2, 3, \dots, N-1\}$

- n-bit address spaces, typical: 32-bit and 64-bit

- ▶ **Physical address space:** Set of $M = 2^m$ physical addresses

$\{0, 1, 2, 3, \dots, M-1\}$

Why Virtual Memory (VM)?

- ▶ Uses main memory efficiently
 - Use *DRAM as a cache* for parts of a virtual address space
- ▶ Simplifies memory management
 - Each *process* gets the *same uniform linear address space*
- ▶ *Isolates address spaces*
 - One process can't interfere with another's memory
 - User program cannot access privileged kernel information and code

VM as a Tool for Caching

- ▶ Conceptually, *virtual memory* is an array of $N (=2^n)$ contiguous bytes stored on disk.
- ▶ The contents of the array on disk are *cached* in *physical memory* (“*DRAM cache*”)
 - These cached blocks are called *pages* (size is $P = 2^p$ bytes)

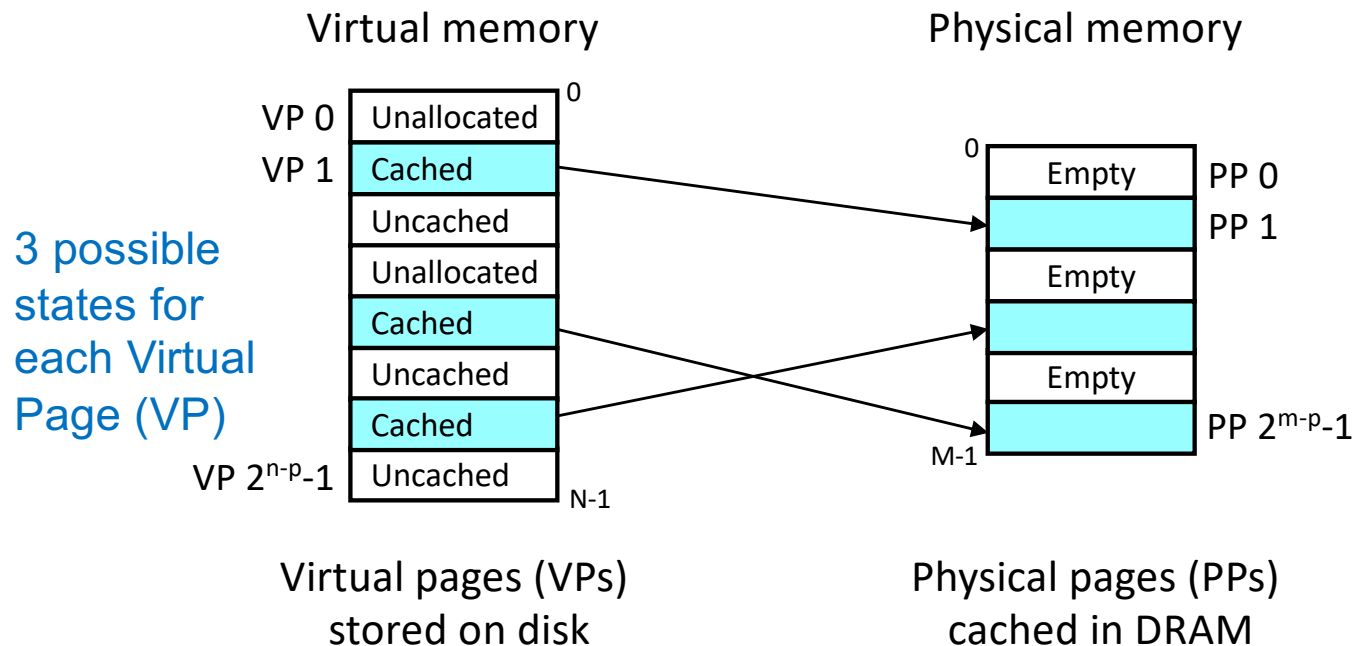


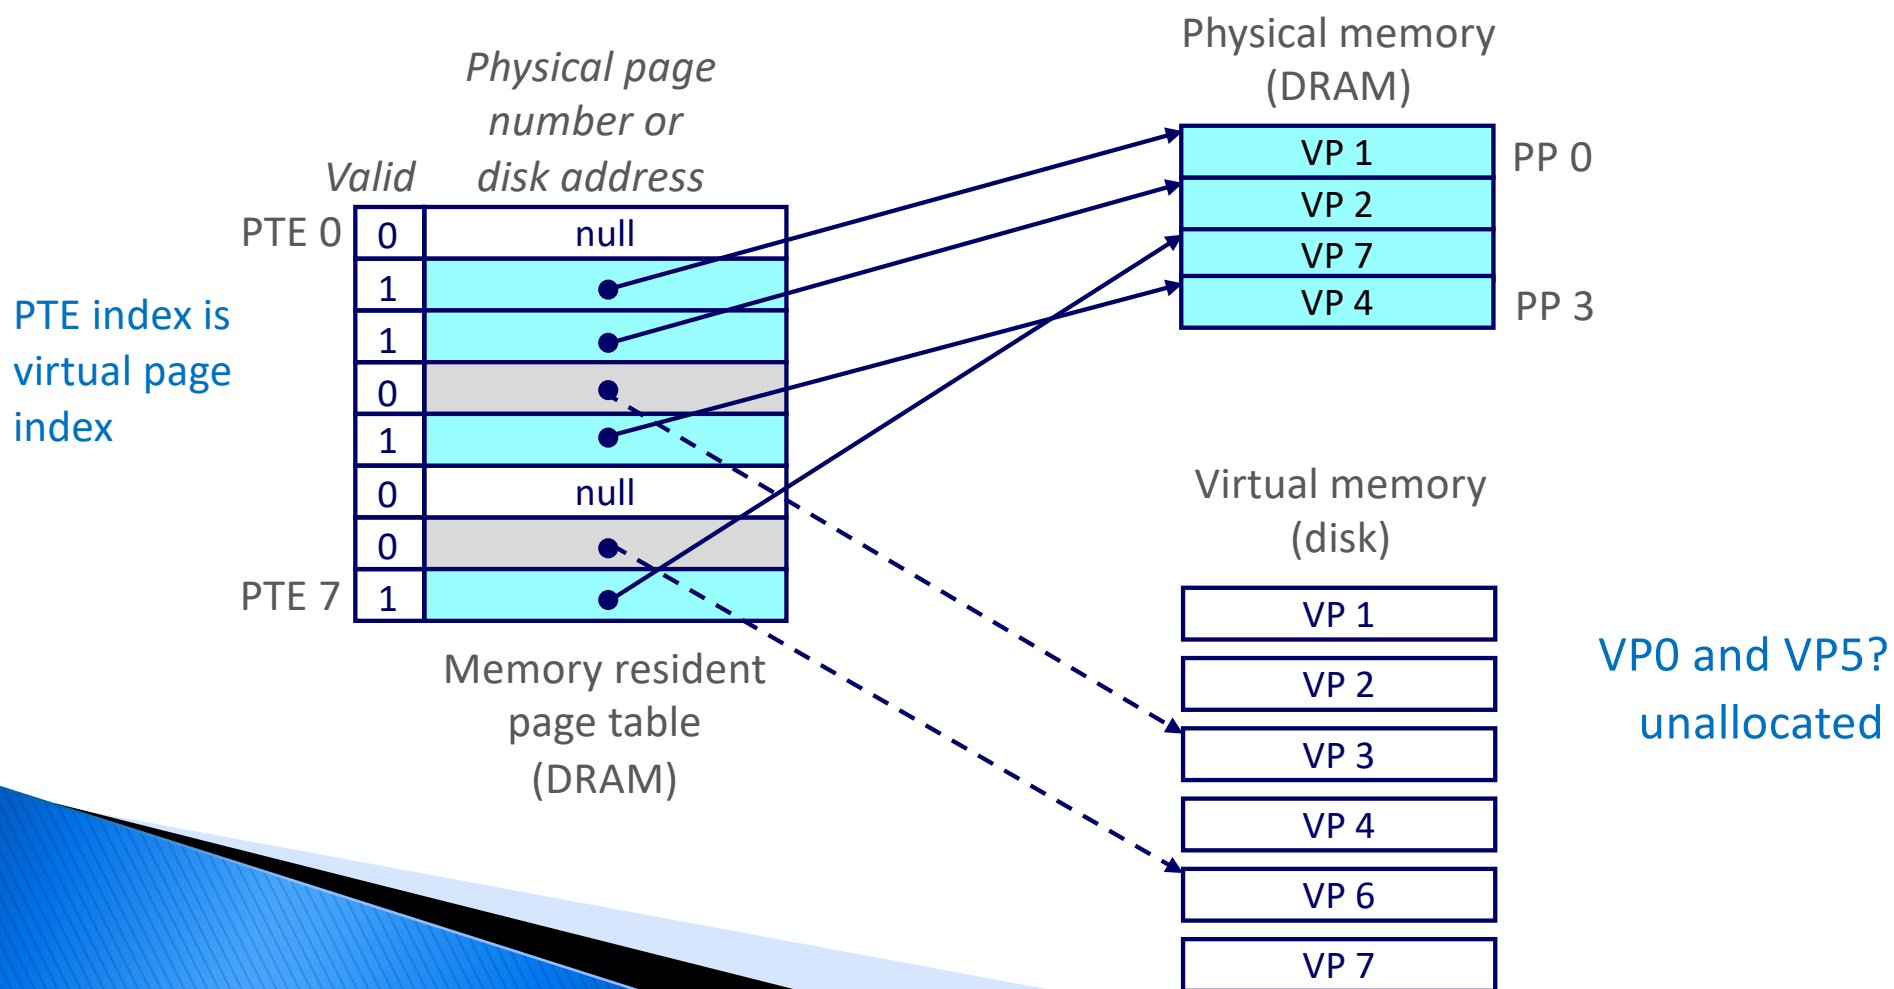
Figure 9.3

“DRAM Cache” Organization (sec 9.3.1)

- ▶ “DRAM cache” (main memory) organization driven by the enormous miss penalty
 - DRAM is about **10x** slower than SRAM (L1, L2, L3 level cache memories)
 - Disk is about **10,000x** slower than DRAM (main memory)
- ▶ Consequences
 - Large page (block) size: typically 4 KB, sometimes 4 MB
 - Fully associative (Recall Chap 6, p. 626, single set holds all the lines)
 - Any VP(virtual page) can be placed in any PP (physical page)
 - Requires a “large” mapping function – different from cache memories
 - Highly sophisticated, expensive replacement algorithms
 - Too complicated and open-ended to be implemented in hardware
 - Write-back (later) rather than write-through

Data Structure: Page Table (enables VM) (sec. 9.3.2)

- ▶ A *page table* is an array of *page table entries (PTEs)* that maps virtual pages to physical pages. State of each VP.
 - One page table (kernel data structure in DRAM) per-process



Exercise

- ▶ Practice problem 9.2 (review terminology + sizes)

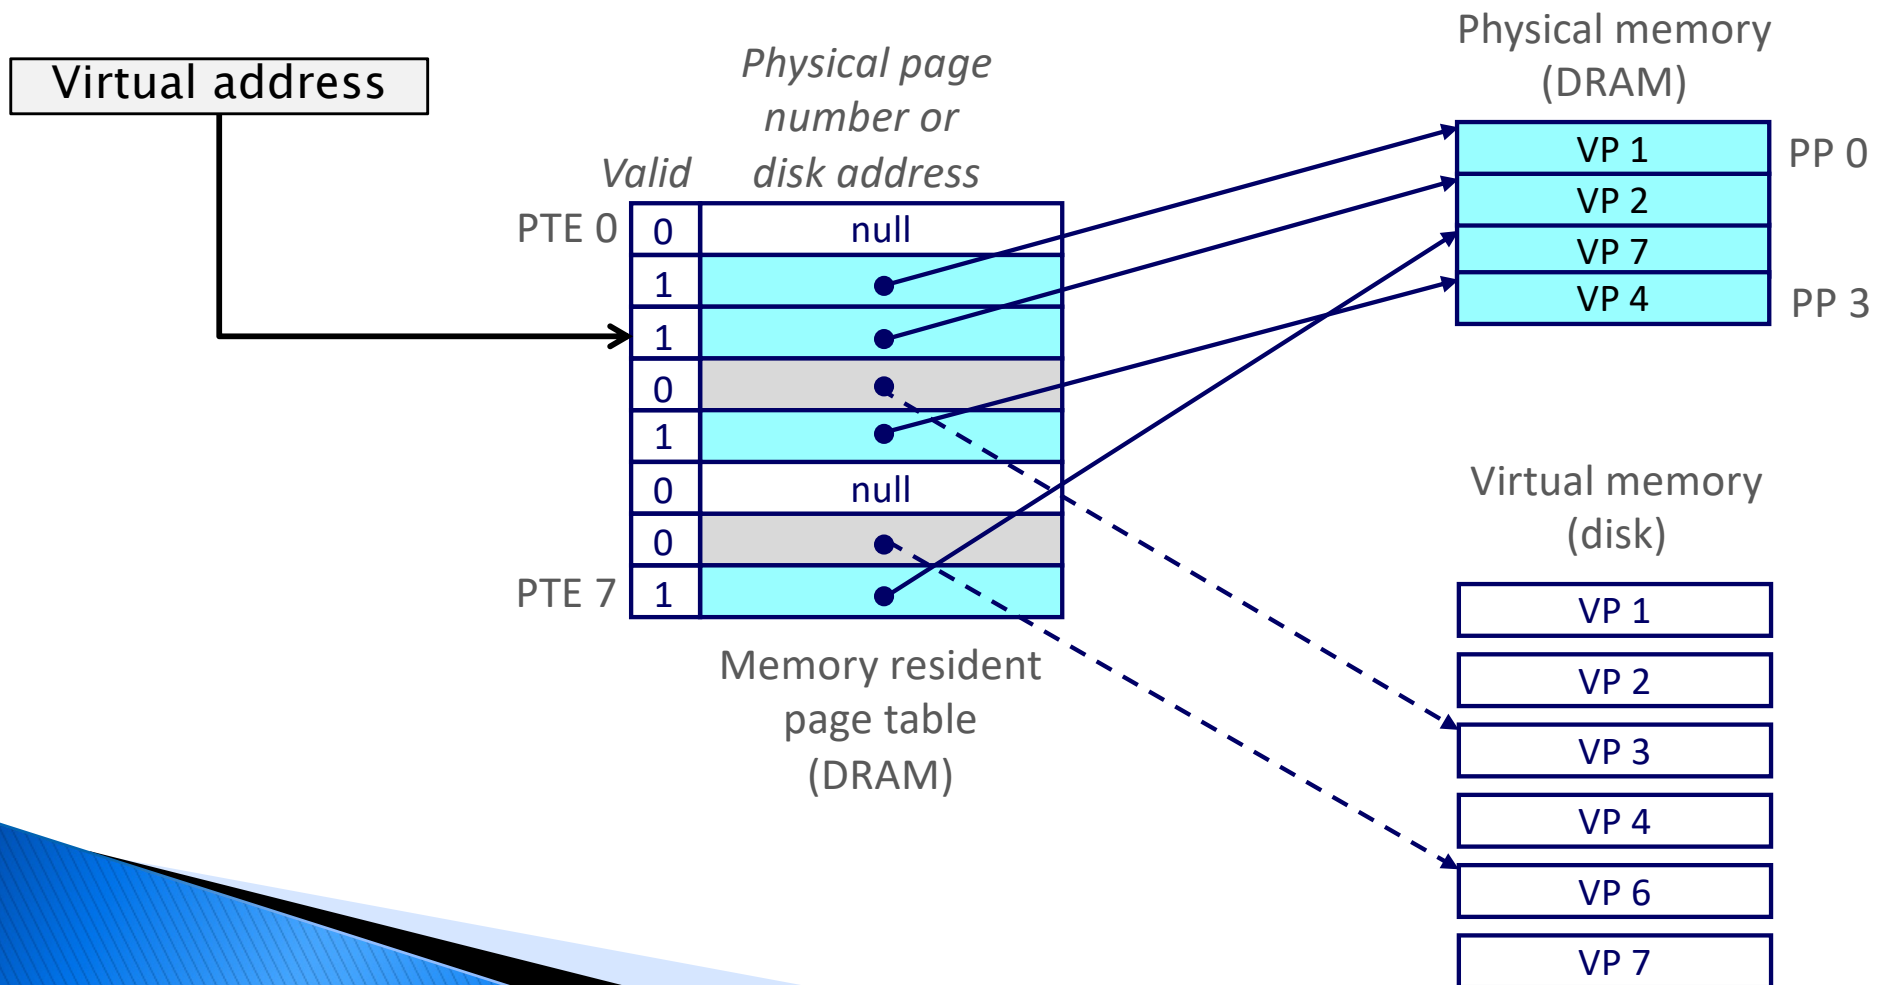
2^n size of virtual memory

n	P (size of page) = 2^p	number of PTE's
16	4K	$2^{16}/2^{12} = 2^4 = 16$ entries
16	8K	$2^{16}/2^{13} = 2^3 = 8$ “
32	4K	$2^{32}/2^{12} = 2^{20} = 1M$ “
32	8K	$2^{32}/2^{13} = 2^{19} = 512K$ ”

$K = 2^{10}$ $M = 2^{20}$ $G = 2^{30}$ $T = 2^{40}$ $Peta = 2^{50}$ $Exa = 2^{60}$

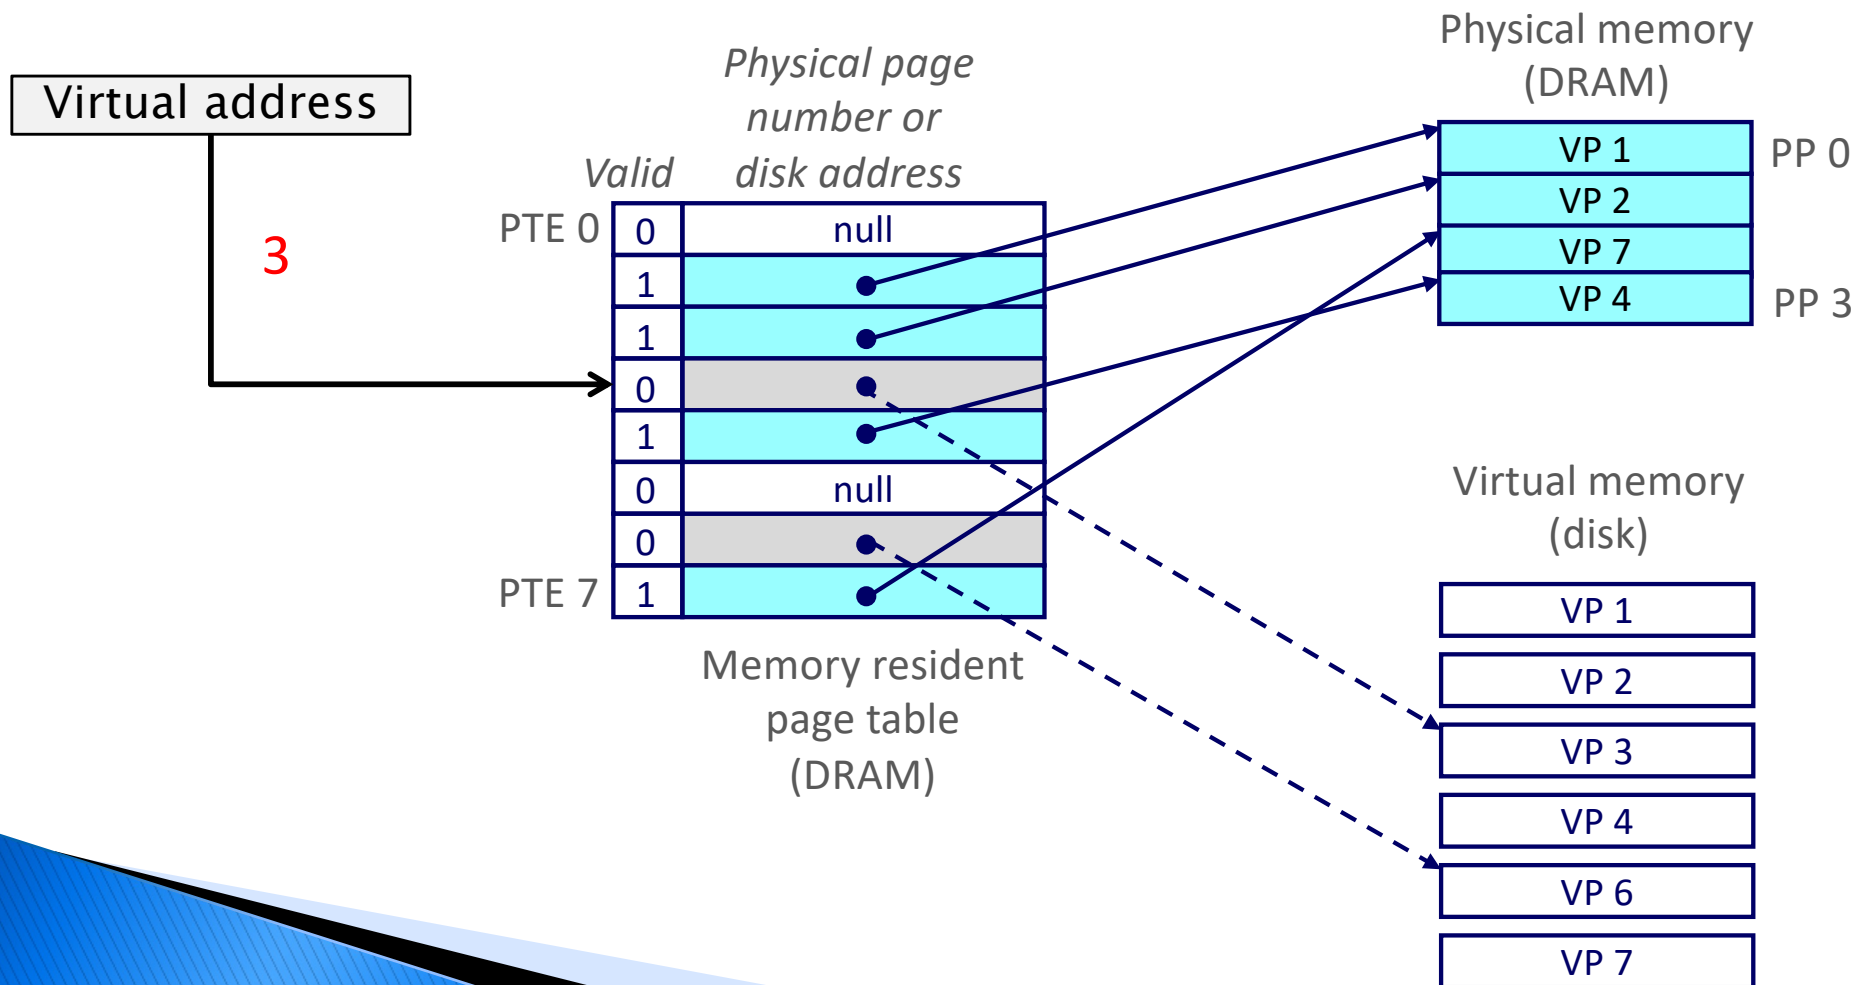
Page Hit (sec. 9.3.3)

- ▶ *Page hit*: reference to *VM word* that is in physical memory (DRAM cache hit)



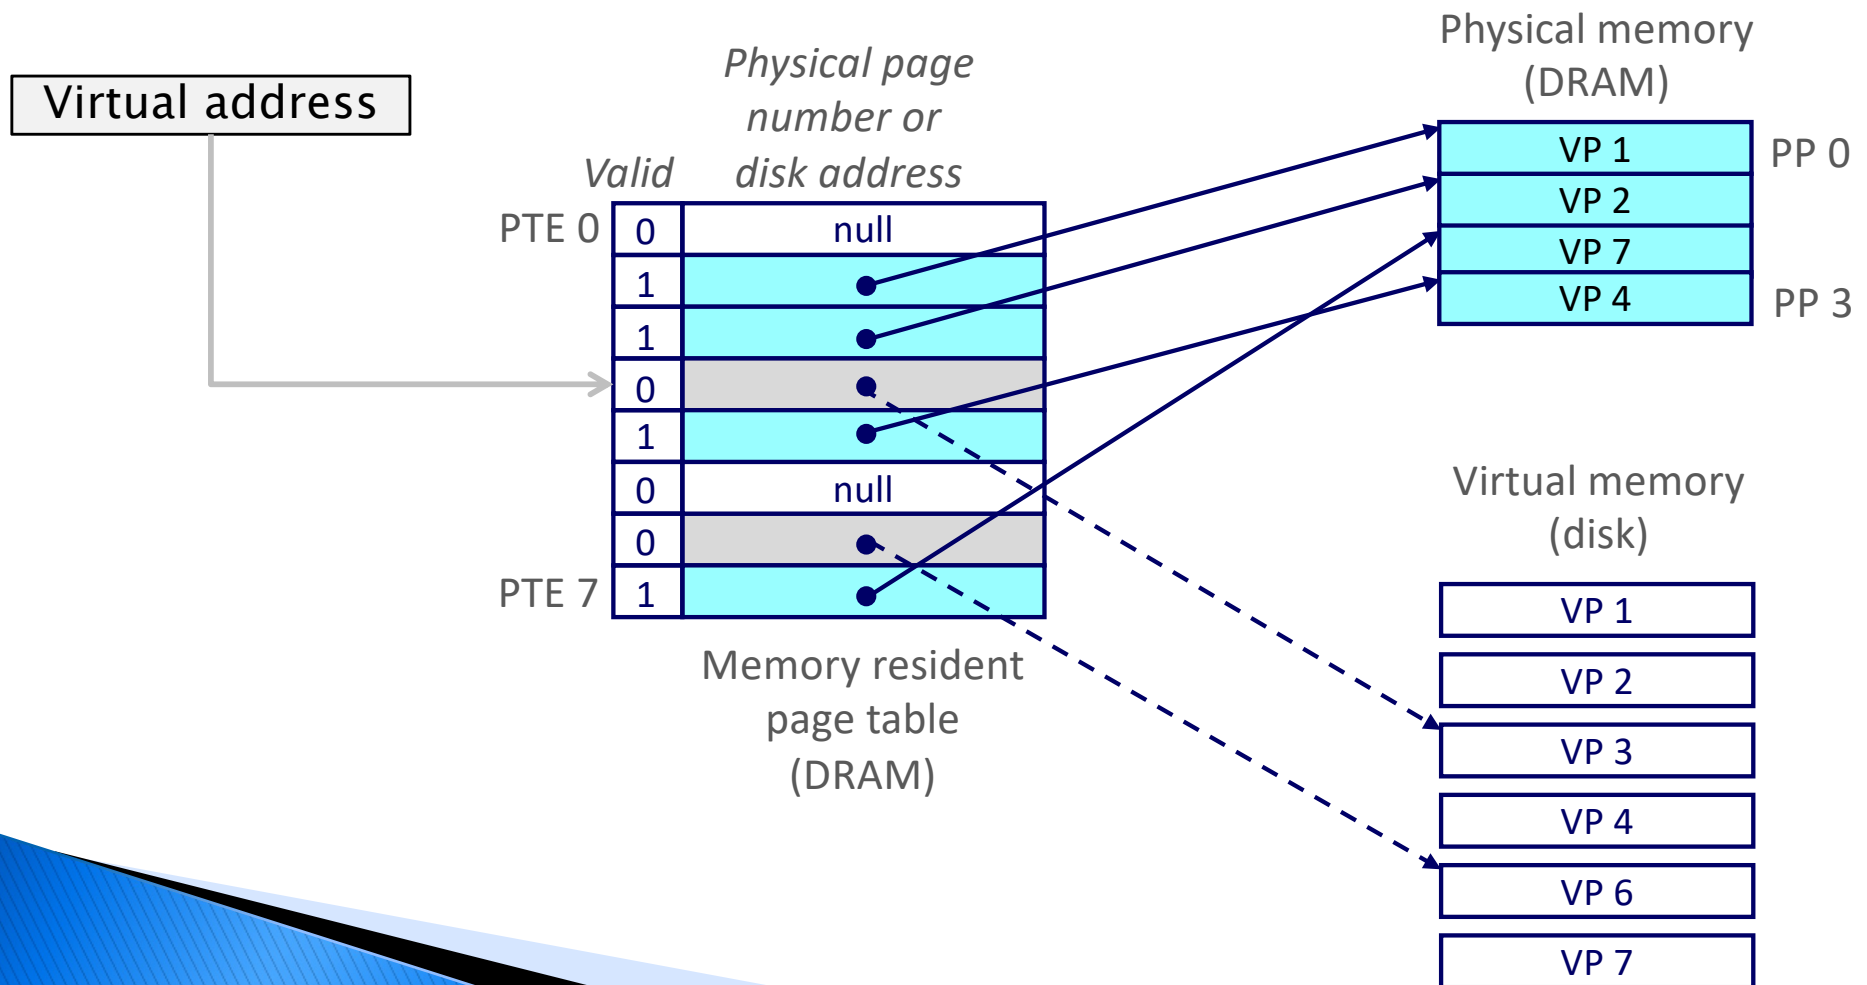
Page Fault (sec. 9.3.4)

- ▶ **Page fault:** reference to VM **page** that is not in physical memory (DRAM cache miss)



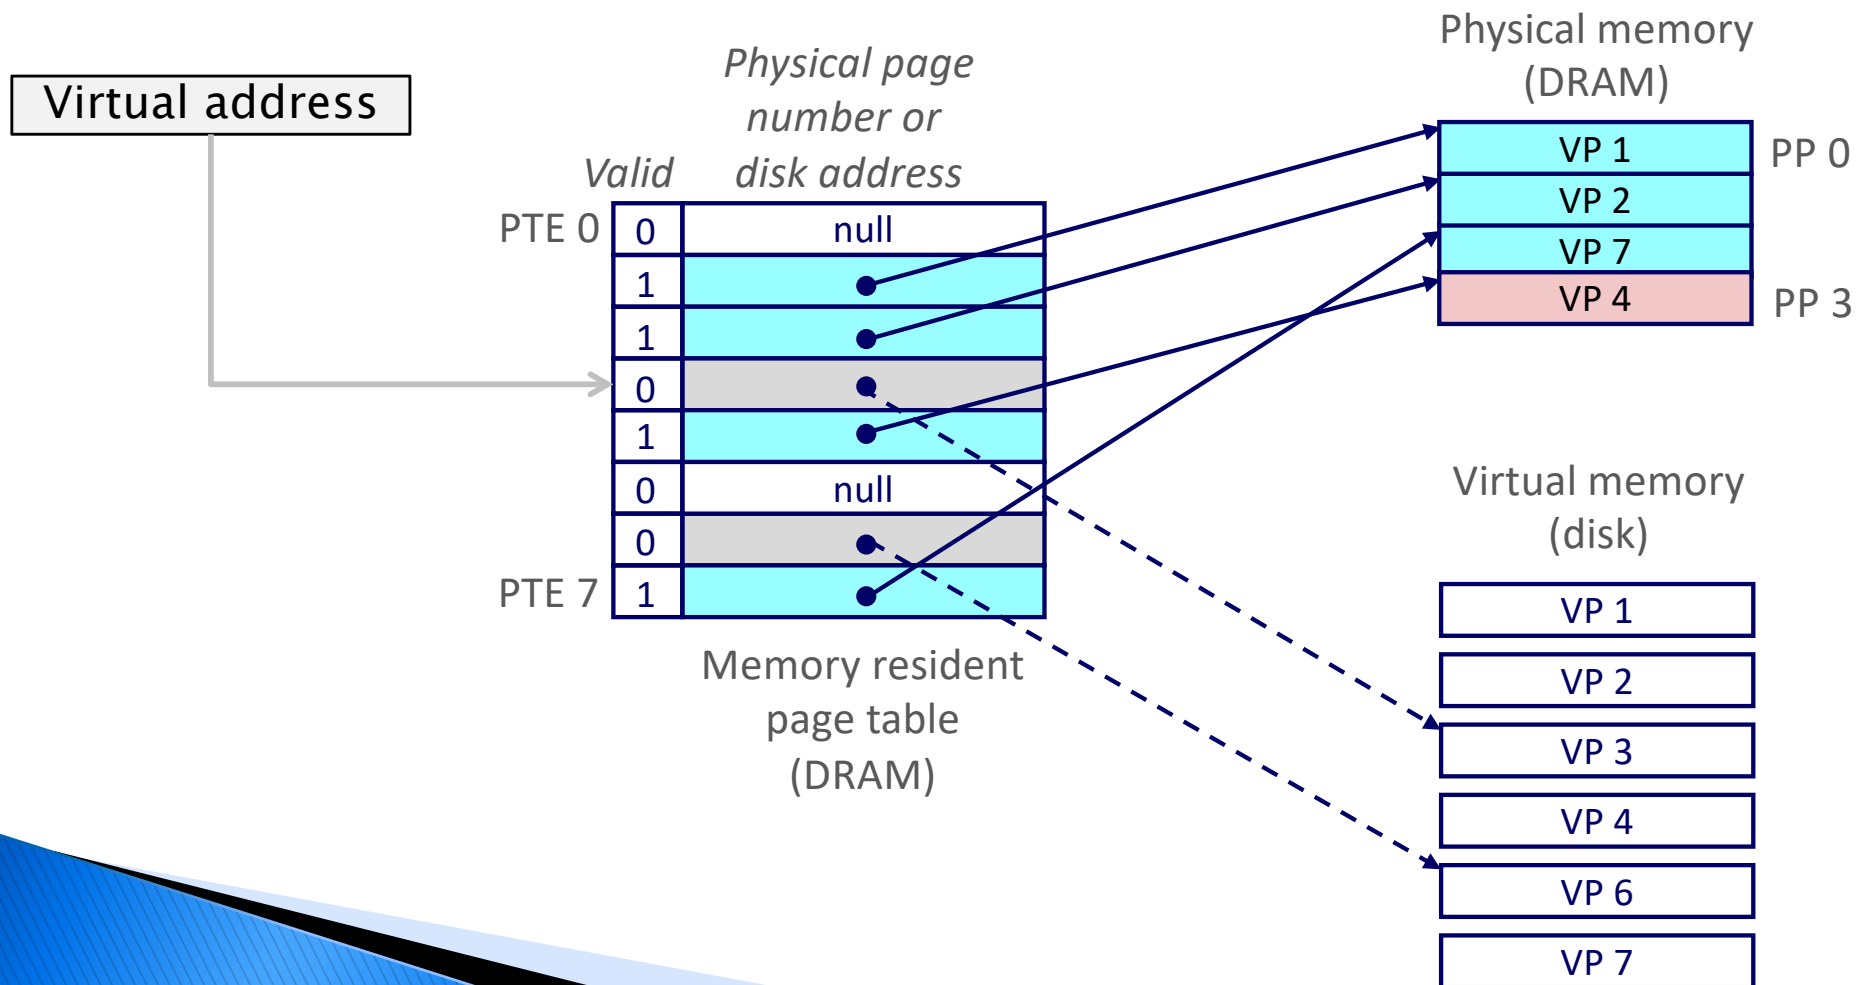
Handling Page Fault

- ▶ **Page miss** causes page fault (an exception)



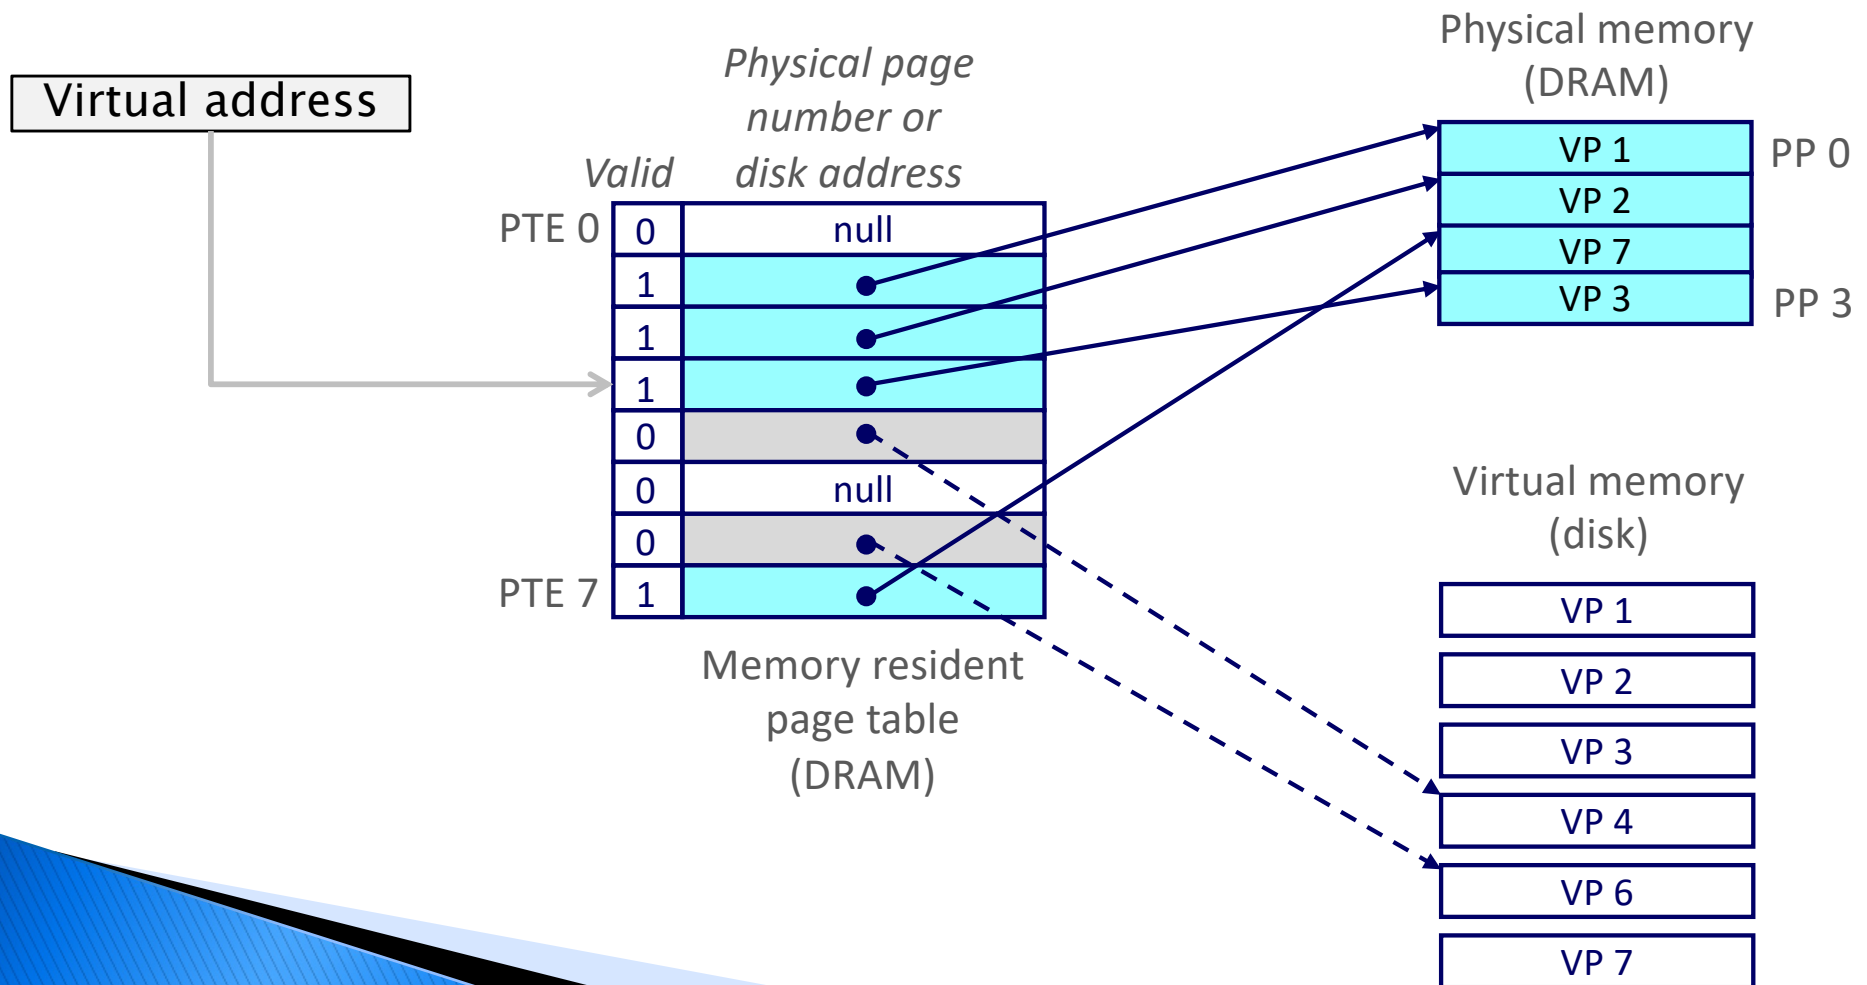
Handling Page Fault

- ▶ Page miss causes page fault (an exception)
- ▶ Page fault handler selects a victim to be evicted (here VP 4)



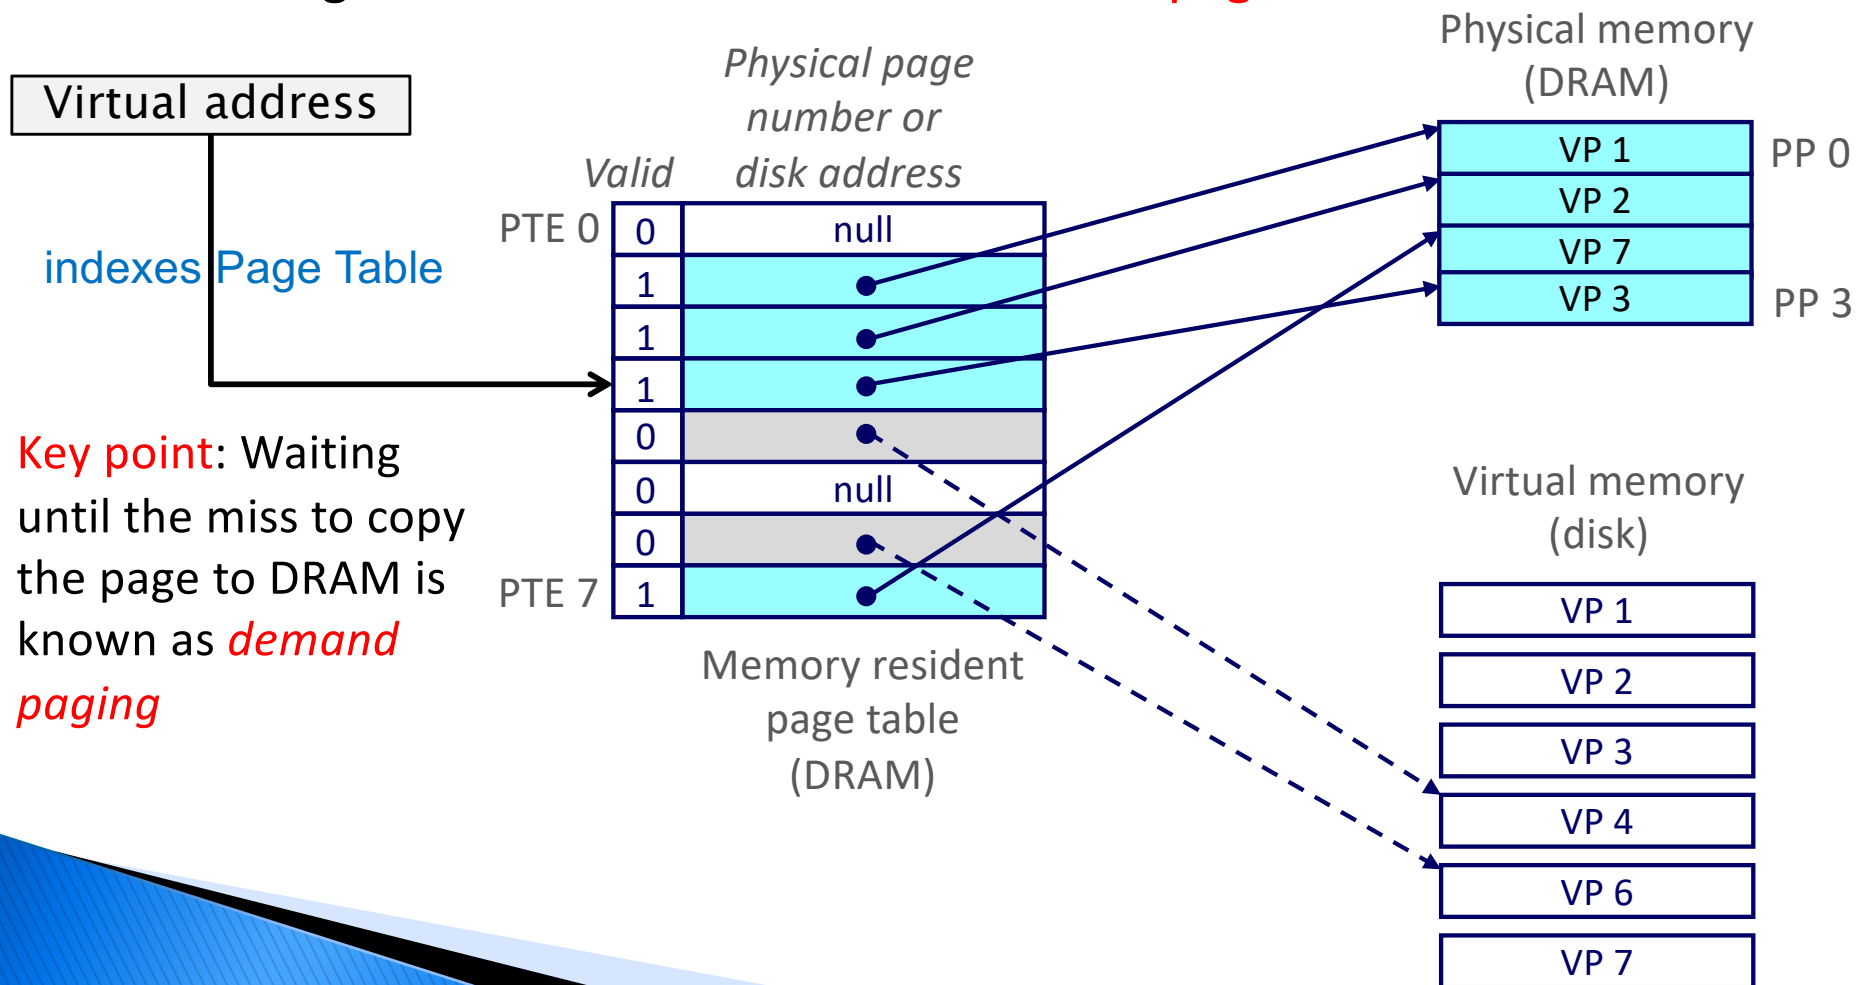
Handling Page Fault

- ▶ Page miss causes page fault (an exception)
- ▶ Page fault handler selects a victim to be evicted (here VP 4)



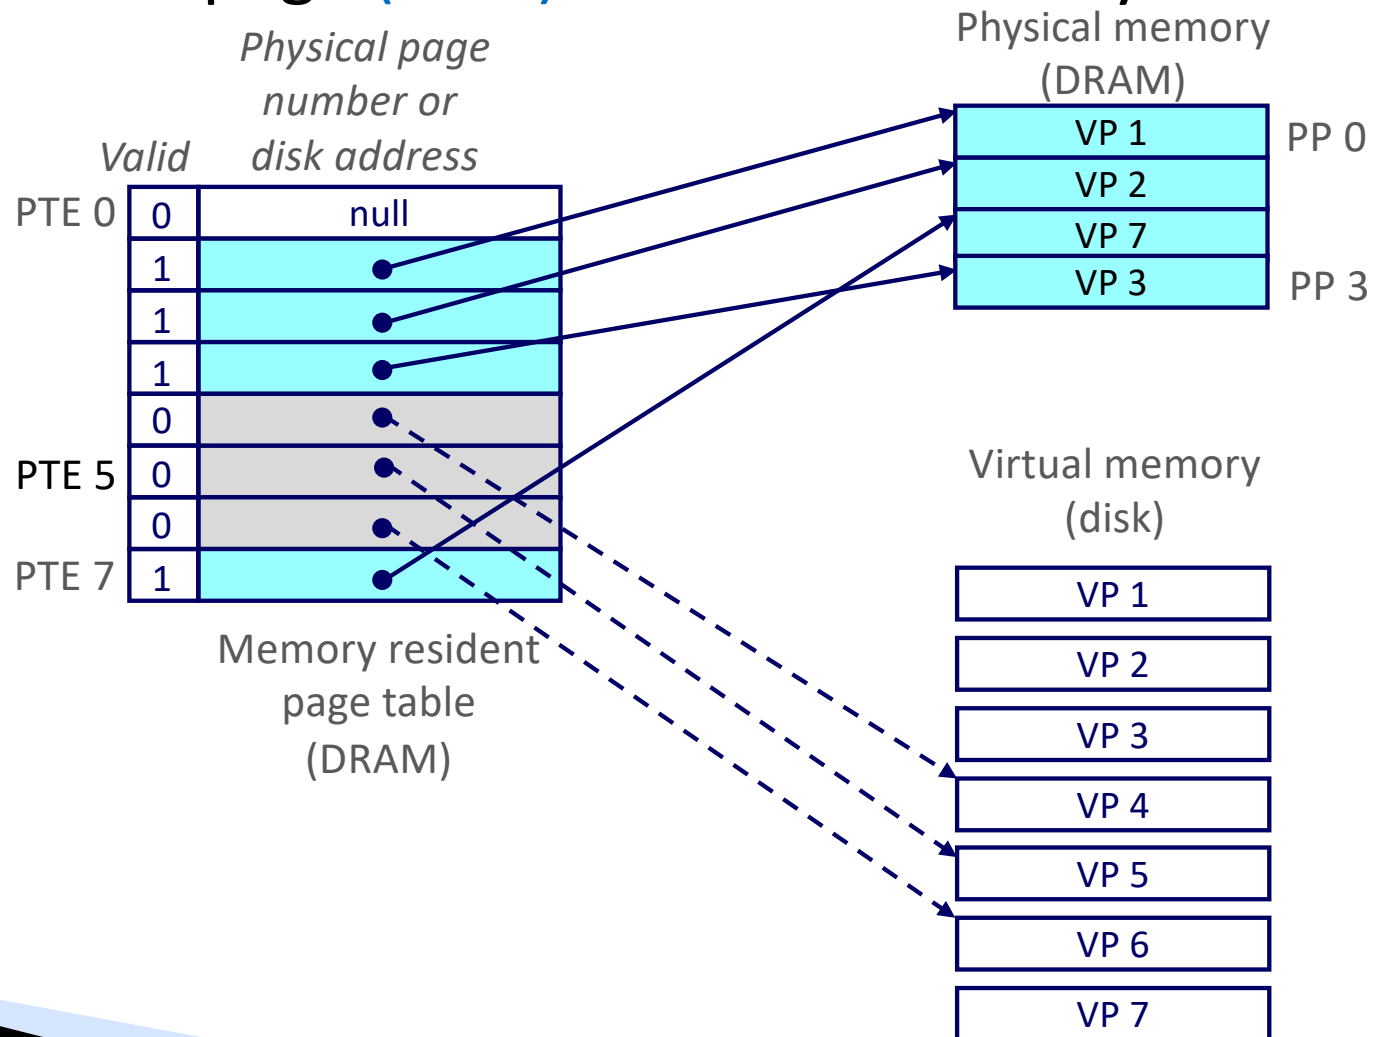
Handling Page Fault

- ▶ Page miss causes page fault (an exception)
- ▶ Page fault handler selects a victim to be evicted (here VP 4) (write back!)
- ▶ Offending instruction is restarted: now it is a **page hit!**



Allocating Pages (sec. 9.3.5 + fig. 9.8)

- ▶ Allocating a new page (VP 5) of virtual memory.



Question (addressed with diagram on the board at the beginning)

- ▶ Given what we know of the cache, what would the process be to provide a Virtual Page to a process that needs data from it?
- ▶ First, recall how was the process to get data from main memory.
- ▶ We will explain it in great detail, but first get a clear picture of what we are trying to do.
- ▶ Keep adding more details to the general picture drawn at the beginning.

Topics

▶ Virtual Memory

- VM as a tool for caching (cont.)
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation
- Simple memory system example

Locality to the Rescue Again! (sec. 9.3.6)

- ▶ Virtual memory seems terribly inefficient, but it works because of locality.
- ▶ At any point in time, programs tend to access a set of active virtual pages called the *working set*
 - Programs with better temporal locality will have smaller working sets and VM works well.
- ▶ If (working set size < main memory size)
 - Good performance for one process after compulsory misses
- ▶ If (SUM(working set sizes) > main memory size)
 - *Thrashing*: Performance meltdown where pages are swapped (copied) in and out continuously