

Lecture #28

System-Level Input/Output (Chap. 10)

Prof. Soraya Abad-Mota, PhD

Today

- ▶ Unix I/O (secs. 10.1 to 10.4)
- ▶ Metadata, sharing, and redirection
- ▶ Standard I/O (part 2)
- ▶ Closing remarks

Input/Output


- ▶ “Process of copying data between main memory and external devices”
- ▶ Examples of external devices: disk drives, terminals, and networks.
- ▶ **Input operation**: copies data from an I/O device to main memory
- ▶ **Output operation**: copies data from memory to a device

Unix I/O Overview

- ▶ A Linux *file* is a sequence of m bytes:
 - $B_0, B_1, \dots, B_k, \dots, B_{m-1}$
- ▶ Cool fact: All I/O devices are represented as files:
 - `/dev/sda2` (`/usr` disk partition)
 - `/dev/tty2` (terminal)
- ▶ Even the kernel is represented as a file:
 - `/boot/vmlinuz-3.13.0-55-generic` (kernel image)
 - `/proc` (kernel data structures)

Unix I/O Overview

- ▶ Elegant mapping of files to devices allows kernel to export simple interface called *Unix I/O*:
 - Opening and closing files
 - **open()** and **close()**
 - Reading and writing a file
 - **read()** and **write()**
 - Changing the **current file position** (seek) (beyond our scope)
 - indicates next offset into file to read or write
 - **lseek()**



Current file position = k

File Types

- ▶ Each file has a *type* indicating its role in the system
 - *Regular file*: Contains arbitrary data
 - *Directory*: Index for a related group of files
 - (more on the next two slides)
 - *Socket*: For communicating with a process on another machine (across a network)
- ▶ Other file types beyond our scope
 - *Named pipes (FIFOs)*
 - *Symbolic links*
 - *Character and block devices*

Regular Files

- ▶ A regular file contains arbitrary data
- ▶ Applications often distinguish between *text files* and *binary files*
 - Text files are regular files with only ASCII or Unicode characters
 - Binary files are everything else
 - e.g., object files, JPEG images
 - Kernel doesn't know the difference!
- ▶ Text file is sequence of *text lines*
 - Text line is sequence of chars terminated by *newline char* ('\n')
 - Newline is 0xa, same as ASCII line feed character (LF)
- ▶ End of line (EOL) indicators in other systems
 - Linux and Mac OS: '\n' (0xa)
 - line feed (LF)
 - Windows and Internet protocols: '\r\n' (0xd 0xa)
 - Carriage return (CR) followed by line feed (LF)




Directories

- ▶ Directory consists of an array of *links*
 - Each link maps a *filename* to a file
- ▶ Each directory contains at least two entries
 - . (dot) is a link to itself
 - .. (dot dot) is a link to *the parent directory* in the *directory hierarchy* (next slide)
 - the command `ls -a` allows you to see both
- ▶ Commands for manipulating directories
 - `mkdir`: create empty directory
 - `ls`: view directory contents
 - `rmdir`: delete empty directory

Directory Hierarchy

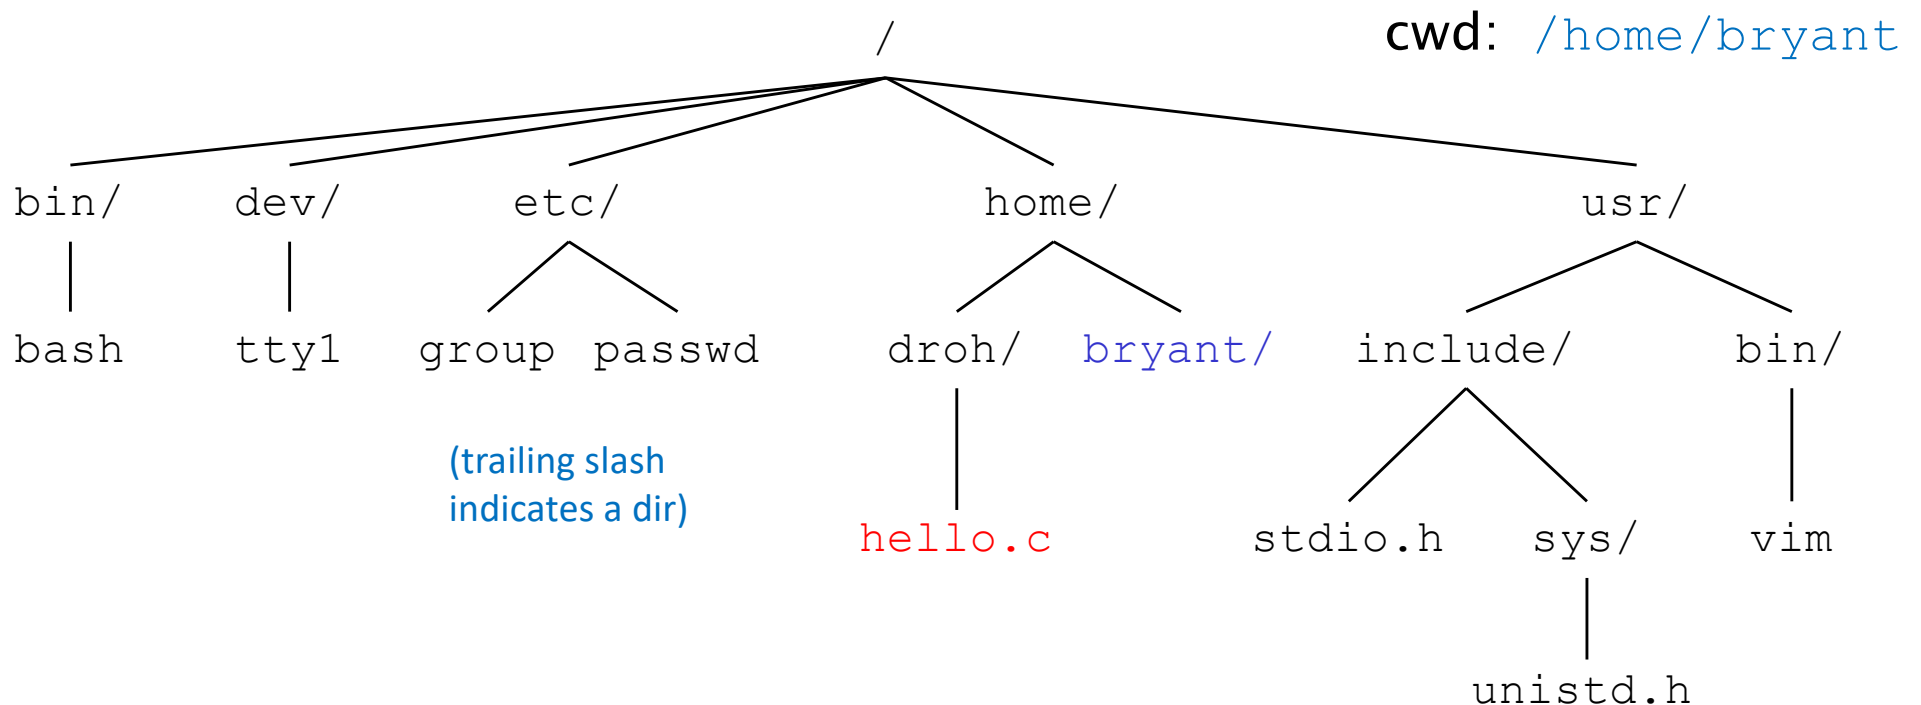
- ▶ All files are organized as a hierarchy anchored by the root directory named / (slash)



- ▶ Kernel maintains *current working directory (cwd)* for each process
 - Modified using the `cd` command; shown by `pwd` command

Pathnames

- ▶ Locations of files in the hierarchy denoted by *pathnames*
 - *Absolute pathname* starts with '/' and denotes path from root
 - /home/droh/hello.c
 - *Relative pathname* denotes path from current working directory
 - ../droh/hello.c



Opening Files

- ▶ Opening a file informs the kernel that you are getting ready to access that file

```
int fd;    /* file descriptor */

if ((fd = open("/etc/hosts", O_RDONLY)) < 0) {
    perror("open");
    exit(1);
}
```

- ▶ Returns a small identifying integer: *file descriptor*
 - **fd == -1** indicates that an error occurred
- ▶ Each process created by a Linux shell begins life with three open files associated with a terminal: (digit is the file descriptor for each)
 - 0: standard input (stdin)
 - 1: standard output (stdout)
 - 2: standard error (stderr)

Closing Files

- ▶ Closing a file informs the kernel that you are finished accessing that file

```
int fd;      /* file descriptor */
int retval;  /* return value */

if ((retval = close(fd)) < 0) {
    perror("close");
    exit(1);
}
```

- ▶ Closing an already closed file is a recipe for disaster in threaded programs (more on this later)
- ▶ Moral: Always check return codes, even for seemingly benign functions such as `close()`

Reading Files

- ▶ Reading a file copies bytes from the current file position to memory, and then updates file position

```
char buf[512];
int fd;          /* file descriptor */
int nbytes;      /* number of bytes read */

/* Open file fd ... */
/* Then read up to 512 bytes from file fd */
if ((nbytes = read(fd, buf, sizeof(buf))) < 0) {
    perror("read");
    exit(1);
}
```

- ▶ Returns number of bytes read from file `fd` into `buf`
 - Return type **`ssize_t`** is signed integer
 - **`nbytes < 0`** indicates that an error occurred
 - **Short counts** (**`nbytes < sizeof(buf)`**) possible and not errors!
 - **`nbytes = 0`** indicates EOF (end of file)

Writing Files

- ▶ Writing a file copies bytes from memory to the current file position, and then updates current file position

```
char buf[512];
int fd;          /* file descriptor */
int nbytes;      /* number of bytes written */

/* Open the file fd ... */
/* Then write up to 512 bytes from buf to file fd */
if ((nbytes = write(fd, buf, sizeof(buf))) < 0) {
    perror("write");
    exit(1);
}
```

- ▶ Returns number of bytes written from `buf` to file `fd`
 - **`nbytes < 0`** indicates that an error occurred
 - As with reads, short counts are possible and are not errors!

Simple Unix I/O example

- ▶ Copying stdin to stdout, one byte at a time

```
#include "csapp.h"

int main(void)
{
    char c;

    while (Read(STDIN_FILENO, &c, 1) != 0)
        Write(STDOUT_FILENO, &c, 1);
    exit(0);
}
```

On Short Counts (pp. 895-897)

- ▶ Short counts can occur in these situations:
 - Encountering (end-of-file) EOF on reads
 - Reading text lines from a terminal
 - Reading and writing network sockets
- ▶ Short counts never occur in these situations:
 - Reading from disk files (except for EOF)
 - Writing to disk files
- ▶ Best practice is to always allow for short counts.

Topics

- ▶ Unix I/O
- ▶ (skipping sec. 10.5 Robust read/write w/ R_{IO} package)
- ▶ Metadata, sharing, and redirection (secs. 10.6 to 10.9)
- ▶ Standard I/O
- ▶ Closing remarks

File Metadata (section 10.6)

- ▶ *Metadata* is data about data, in this case file data
- ▶ Per-file metadata maintained by kernel: accessed by users with the **stat**(filename) and **fstat**(fd) functions

```
/* Metadata returned by the stat and fstat functions */
struct stat {
    dev_t      st_dev;      /* Device */
    ino_t      st_ino;      /* inode */
    mode_t     st_mode;     /* Protection and file type */
    nlink_t    st_nlink;    /* Number of hard links */
    uid_t      st_uid;      /* User ID of owner */
    gid_t      st_gid;      /* Group ID of owner */
    dev_t      st_rdev;     /* Device type (if inode device) */
    off_t      st_size;     /* Total size, in bytes */
    unsigned long st_blksize; /* Blocksize for filesystem I/O */
    unsigned long st_blocks; /* Number of blocks allocated */
    time_t     st_atime;    /* Time of last access */
    time_t     st_mtime;    /* Time of last modification */
    time_t     st_ctime;    /* Time of last change */
};
```

Example of Accessing File Metadata

```
int main (int argc, char **argv)
{
    struct stat stat;
    char *type, *readok;

    Stat(argv[1], &stat);
    if (S_ISREG(stat.st_mode)) /* Determine file type */
        type = "regular";
    else if (S_ISDIR(stat.st_mode))
        type = "directory";
    else
        type = "other";
    if ((stat.st_mode & S_IRUSR)) /* Check read access */
        readok = "yes";
    else
        readok = "no";

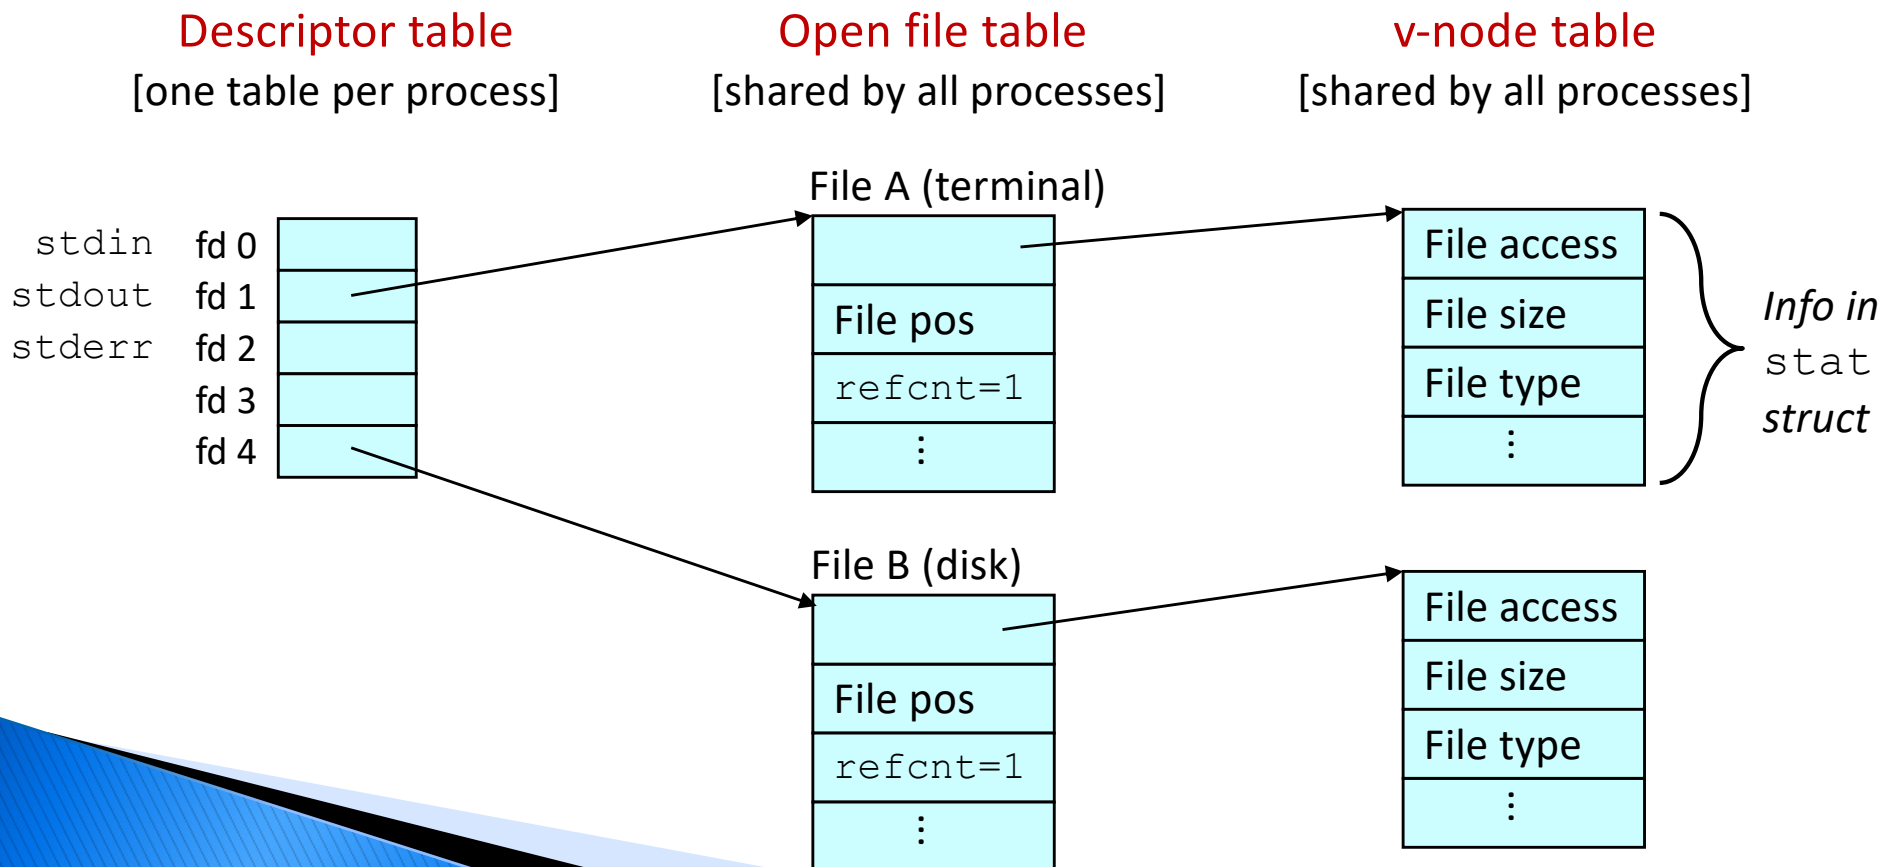
    printf("type: %s, read: %s\n", type, readok);
    exit(0);
}
```

statcheck.c

```
linux> ./statcheck statcheck.c
type: regular, read: yes
linux> chmod 000 statcheck.c
linux> ./statcheck statcheck.c
type: regular, read: no
linux> ./statcheck ..
type: directory, read: yes
```

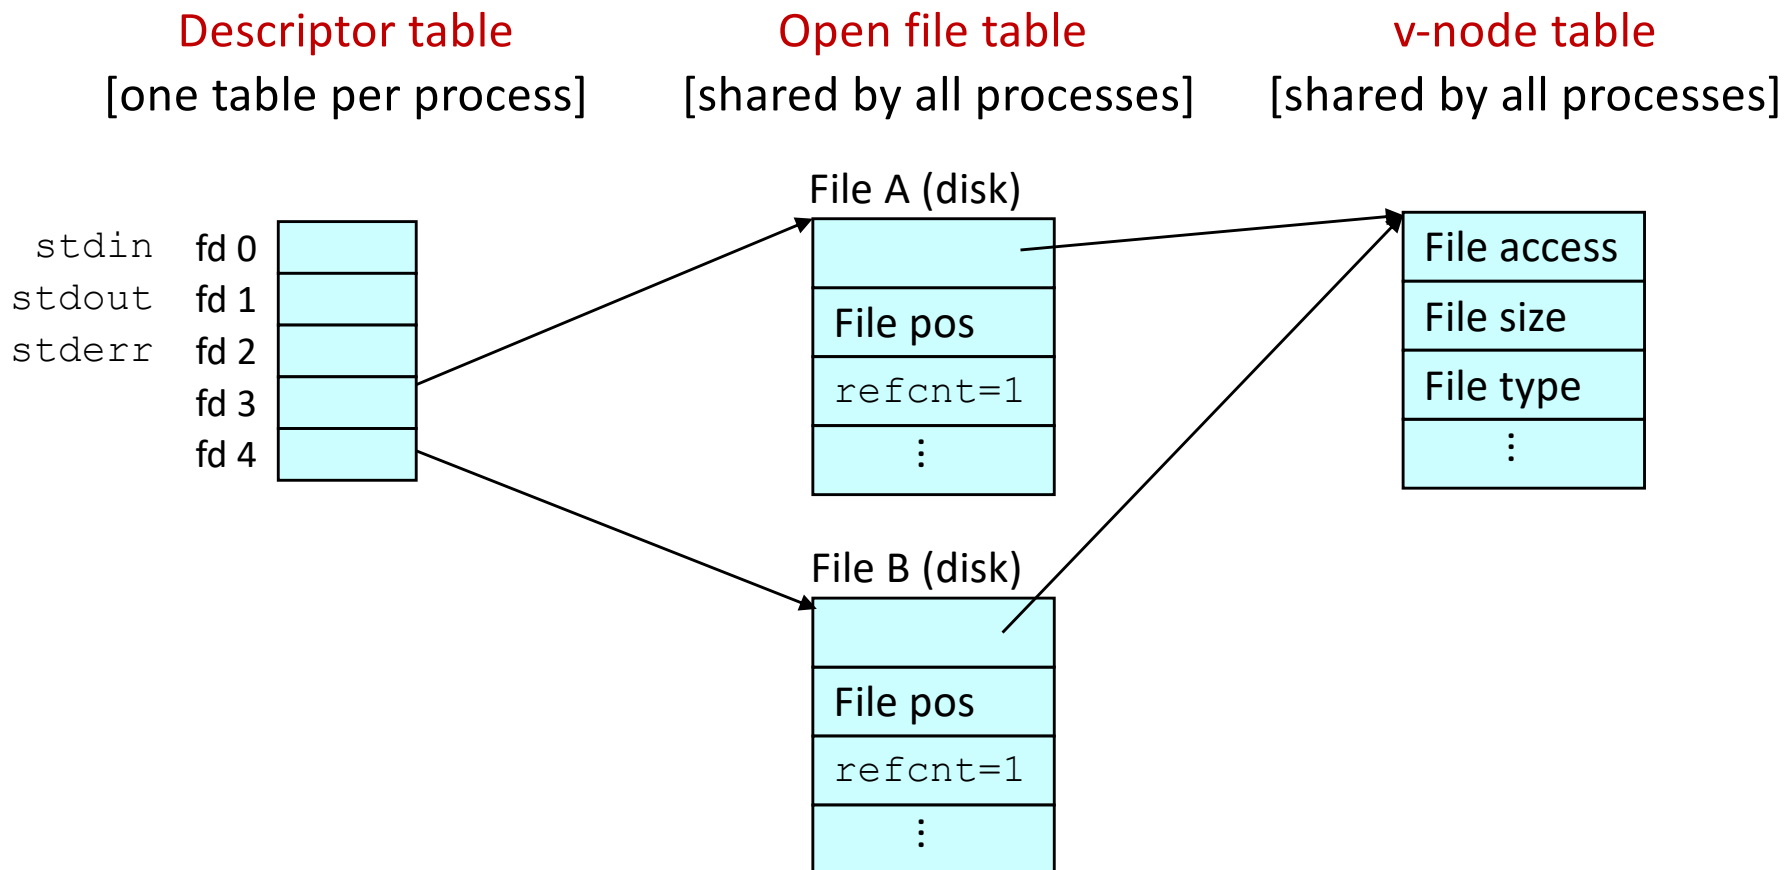
How the Unix Kernel Represents Open Files

- Two descriptors referencing two distinct open files.
Descriptor 1 (stdout) points to terminal, and descriptor 4 points to open disk file (fig. 10.12)



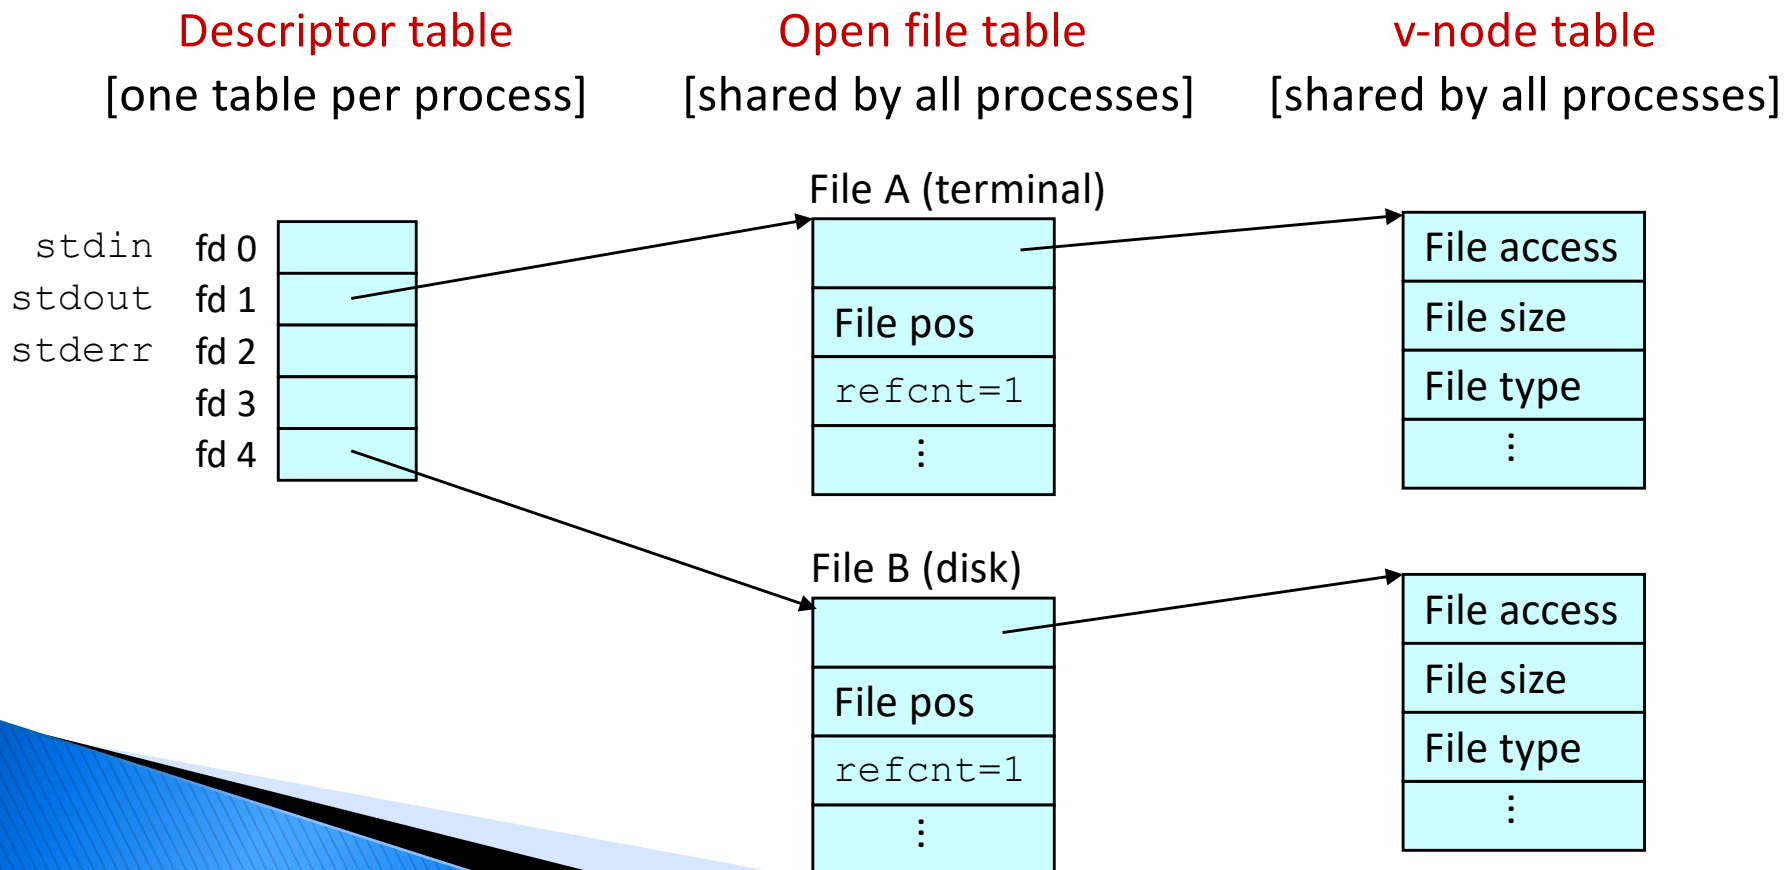
File Sharing

- ▶ Two distinct descriptors sharing the same disk file through two distinct open file table entries (fig. 10.13)
 - E.g. Calling **open** twice with the same **filename** argument



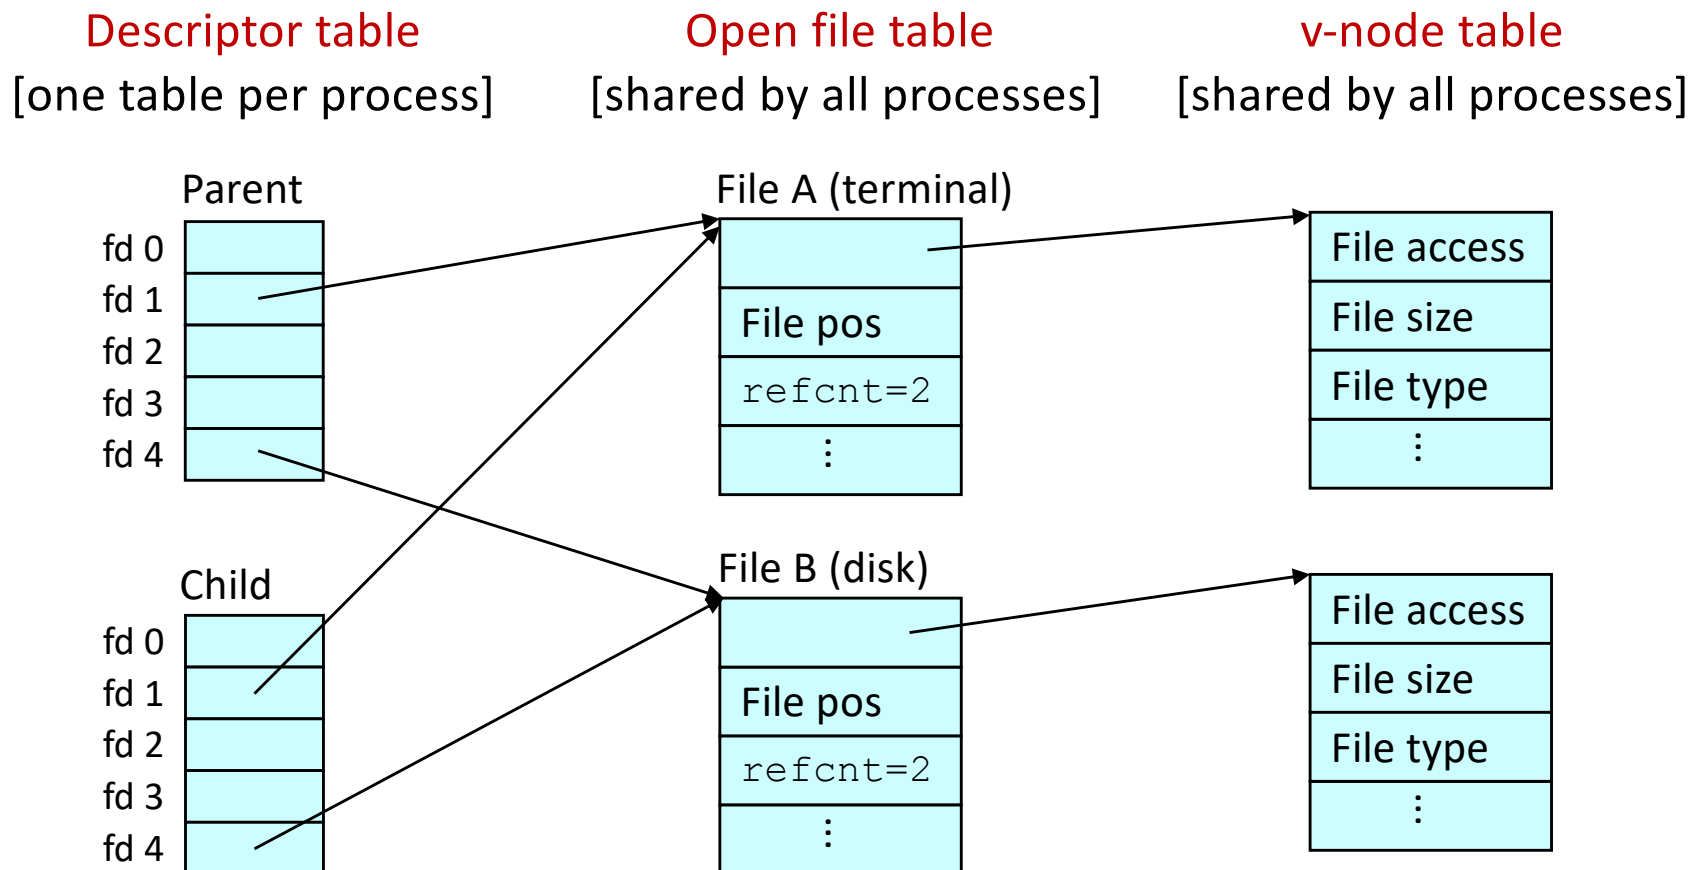
How Processes Share Files: `fork`

- ▶ A child process inherits its parent's open files
 - Note: situation unchanged by **`exec`** functions (use **`fcntl`** to change)
- ▶ ***Before*** `fork` call: (fig. 10.12)



How Processes Share Files: `fork` (2)

- ▶ A child process inherits its parent's open files
- ▶ *After* `fork`: (fig. 10.14)
 - Child's table same as parent's, and +1 to each refcnt



I/O Redirection

- ▶ Question: How does a shell implement I/O redirection?

linux> ls > foo.txt

- ▶ Answer: By calling the `dup2 (oldfd, newfd)` function
 - Copies (per-process) descriptor table entry **oldfd** to entry **newfd**

Descriptor table

before `dup2 (4, 1)`

fd 0	
fd 1	a
fd 2	
fd 3	
fd 4	b



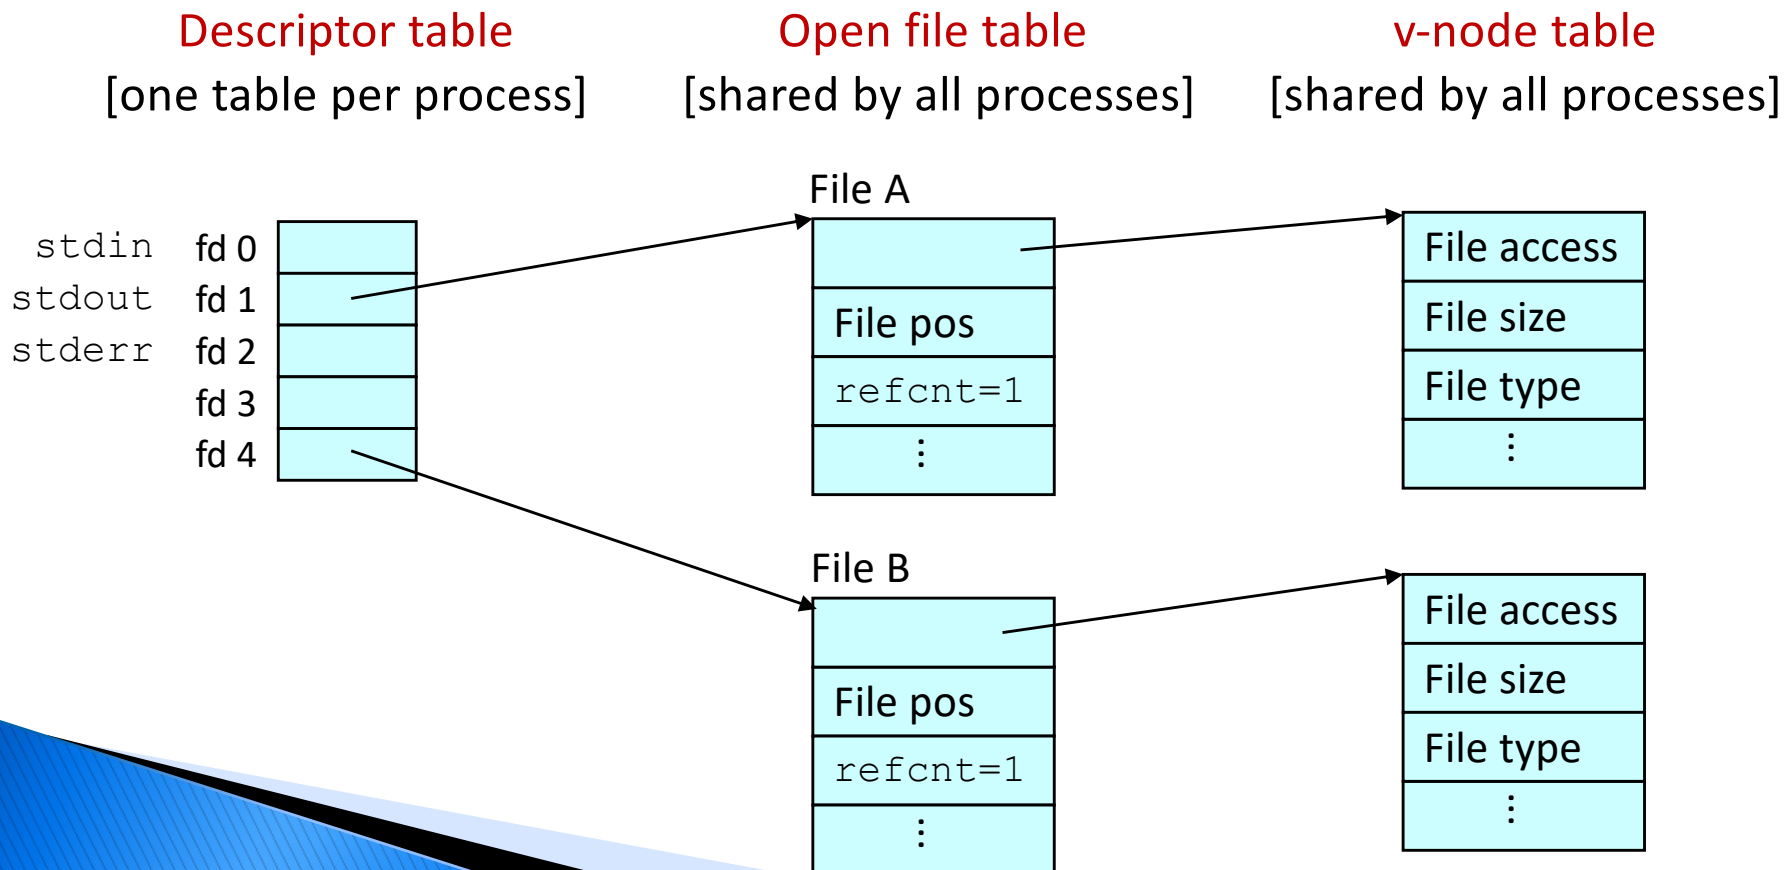
Descriptor table

after `dup2 (4, 1)`

fd 0	
fd 1	b
fd 2	
fd 3	
fd 4	b

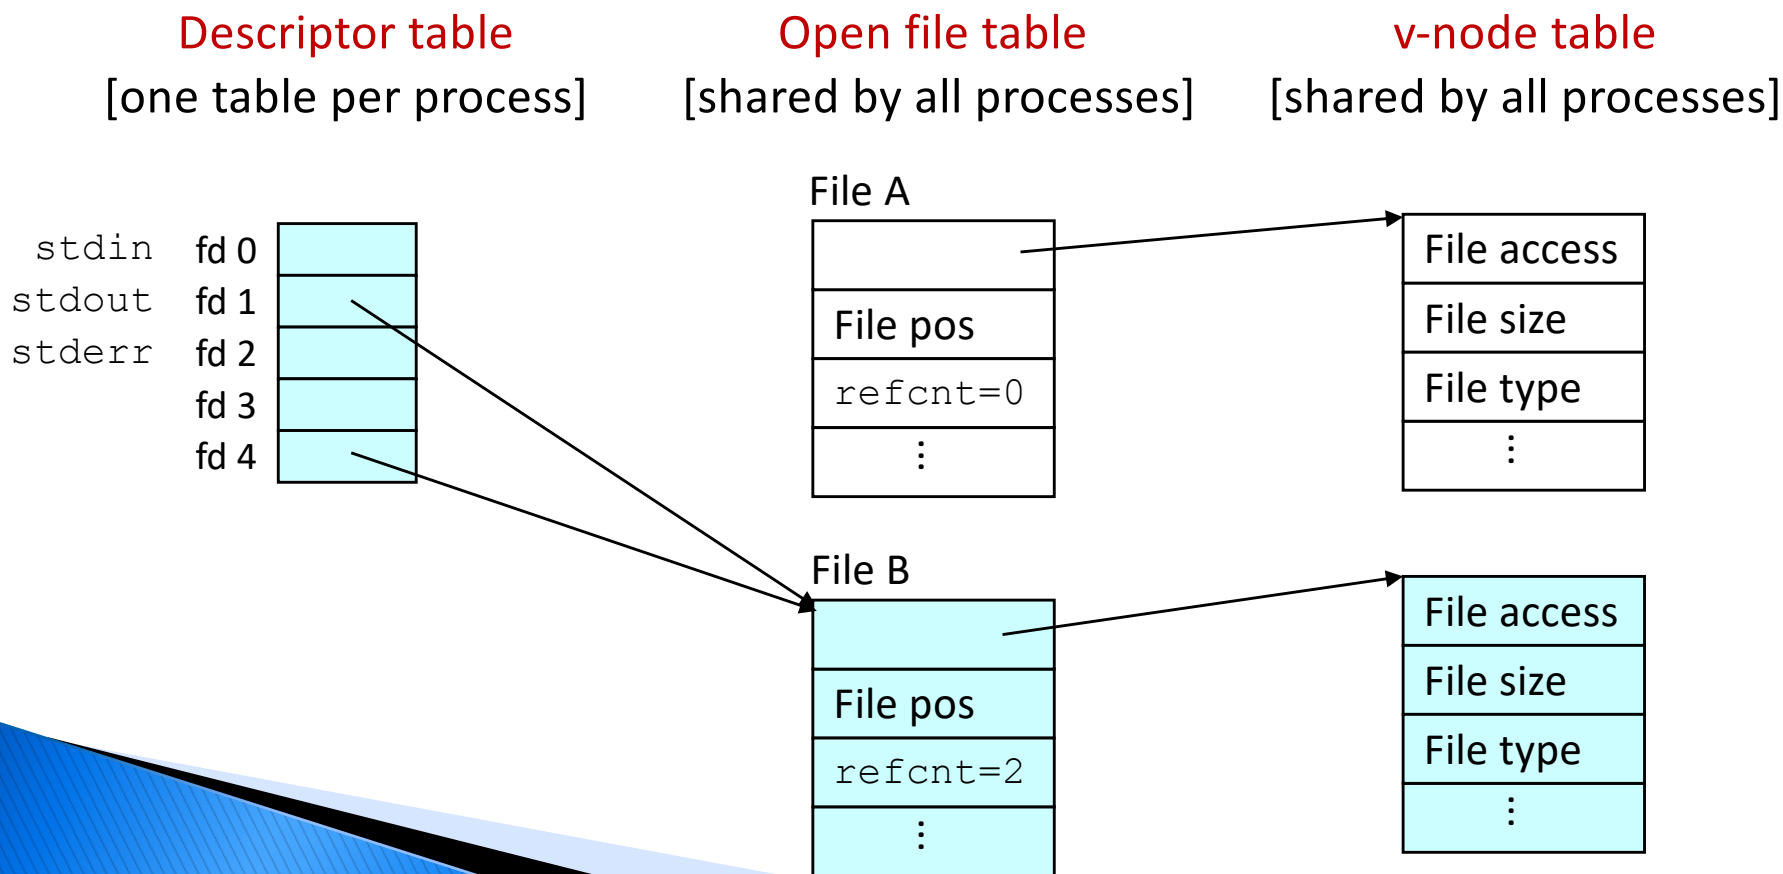
I/O Redirection Example

- ▶ Step #1: open file to which stdout should be redirected
 - Happens in child executing shell code, before **exec**



I/O Redirection Example (cont.)

- ▶ Step #2: call `dup2 (4, 1)`
 - cause fd=1 (stdout) to refer to disk file pointed at by fd=4



Two exercises

- ▶ Problem 10.2
- ▶ Problem 10.3 (careful with the answer given in the book, there is a subtlety here)
- ▶ Both use file foobar.txt which contains:

foobar

Only those 6 characters and the eof char

Problem 10.2 (with minor changes)

```
#include "csapp.h"

int main()
{
    int fd1, fd2;
    char c;

    fd1 = open("foobar.txt", O_RDONLY, 0);
    fd2 = open("foobar.txt", O_RDONLY, 0);
    read(fd1, &c, 1);
    printf("c1 = %c\n", c);
    read(fd2, &c, 1);
    printf("c1 = %c\n", c);
    exit(0);
}
```

What is the output?
Build the 3 data
structures.

Problem 10.3

```
#include "csapp.h"
int main()
{
    int fd;
    char c;
    fd1 = Open("foobar.txt", O_RDONLY, 0);
    if (Fork() == 0) {
        Read(fd, &c, 1);
        exit(0);
    }
    Wait(NULL);
    Read(fd, &c, 1);
    printf("c1 = %c\n", c);
    exit(0);
}
```

What is the output?
Build the 3 data
structures.

Topics

- ▶ Unix I/O
- ▶ (skipping sec. 10.5 Robust read/write w/ R_{IO} package)
- ▶ Metadata, sharing, and redirection (secs. 10.6 to 10.9)
- ▶ **Standard I/O**
- ▶ Closing remarks

Standard I/O Functions

- ▶ The C standard library (`libc.so`) contains a collection of higher-level *standard I/O* functions
 - Documented in Appendix B of K&R [[KRC book](#)]
- ▶ Examples of standard I/O functions:
 - Opening and closing files (**fopen** and **fclose**)
 - Reading and writing bytes (**fread** and **fwrite**)
 - Reading and writing text lines (**fgets** and **fputs**)
 - Formatted reading and writing (**fscanf** and **fprintf**)

Standard I/O Streams


- ▶ Standard I/O models open files as *streams*
 - Abstraction for a file descriptor and a buffer in memory
- ▶ C programs “*begin life*” with three open streams (defined in `stdio.h`)
 - **stdin** (standard input)
 - **stdout** (standard output)
 - **stderr** (standard error)

```
#include <stdio.h>
extern FILE *stdin; /* standard input (descriptor 0) */
extern FILE *stdout; /* standard output (descriptor 1) */
extern FILE *stderr; /* standard error (descriptor 2) */

int main() {
    fprintf(stdout, "Hello, world\n");
}
```



Buffered I/O: Motivation

- ▶ Applications often read/write one character at a time
 - `getc`, `putc`, `ungetc`
 - `gets`, `fgets`
 - Read line of text one character at a time, stopping at newline
- ▶ Implementing as Unix I/O calls expensive
 - `read` and `write` require Unix kernel calls
 - > 10,000 clock cycles
- ▶ Solution: Buffered read
 - Use Unix `read` to grab block of bytes
 - User input functions take one byte at a time from buffer
 - Refill buffer when empty



Buffering in Standard I/O

- ▶ Standard I/O functions use buffered I/O



- ▶ Buffer flushed to output fd on “\n”, call to `fflush` or `exit`, or return from `main`.

Standard I/O Buffering in Action

- ▶ You can see this buffering in action for yourself, using the always fascinating Linux `strace` program:

```
#include <stdio.h>

int main()
{
    printf("h");
    printf("e");
    printf("l");
    printf("l");
    printf("o");
    printf("\n");
    fflush(stdout);
    exit(0);
}
```

```
linux> strace ./hello
execve("./hello", ["hello"], [/* ... */]).
...
write(1, "hello\n", 6)                = 6
...
exit_group(0)                         = ?
```

Topics

- ▶ Unix I/O
- ▶ (skipping sec. 10.5 Robust read/write w/ R_{IO} package)
- ▶ Metadata, sharing, and redirection (secs. 10.6 to 10.9)
- ▶ Standard I/O
- ▶ Closing remarks

Pros and Cons of Unix I/O (compared to std I/O)

▶ Pros

- Unix I/O is the most general and lowest overhead form of I/O
 - All other I/O packages are implemented using Unix I/O functions
- Unix I/O provides functions for accessing file metadata
- Unix I/O functions are async-signal-safe and can be used safely in signal handlers

▶ Cons

- Dealing with short counts is tricky and error prone
- Efficient reading of text lines requires some form of buffering, also tricky and error prone
- Both of these issues are addressed by the standard I/O packages

Pros and Cons of Standard I/O (over Unix I/O)

▶ Pros:

- Buffering increases efficiency by decreasing the number of **read** and **write** system calls
- Short counts are handled automatically

▶ Cons:

- Provides no function for accessing file metadata
- Standard I/O functions are not async-signal-safe, and not appropriate for signal handlers
- Standard I/O is not appropriate for input and output on network sockets
 - There are poorly documented restrictions on streams that interact badly with restrictions on sockets (CS:APP3e, Sec 10.11)

Choosing I/O Functions

- ▶ General rule: use the highest-level I/O functions you can
 - Many C programmers are able to do all of their work using the standard I/O functions
 - But, be sure to understand the functions you use!
- ▶ When to use standard I/O
 - When working with disk or terminal files
- ▶ When to use raw Unix I/O
 - Inside signal handlers, because Unix I/O is async-signal-safe
 - In rare cases when you need absolute highest performance

Aside: Working with Binary Files

- ▶ Functions you should never use on binary files
 - Text-oriented I/O such as **fgets**, **scanf**
 - Interpret EOL characters.
 - String functions
 - **strlen**, **strcpy**, **strcat**
 - Interprets byte value 0 (end of string) as special

For Further Information (important resources)

▶ The Unix bible:

- W. Richard Stevens & Stephen A. Rago, ***Advanced Programming in the Unix Environment***, 3rd Edition, Pearson Education Inc., 2013
 - Updated from Stevens's 1993 classic text and from 2nd Ed. 2005
- <http://www.apuebook.com/code3e.html> has code and relevant info about the 3rd ed. of this book

▶ The Linux bible:

- Michael Kerrisk, *The Linux Programming Interface*, No Starch Press, 2010
 - Encyclopedic and authoritative

Continue with last topic

- ▶ Ch. 12: Concurrency (next pptx file)